# The Demographic Tide: A Complete Guide to Understanding How Retirement Flows Will Reshape Equity Markets

## A Teaching Companion for the BDBA Capstone Project

*Written for Antonio — IE University, February 2026*

---

# Table of Contents

## Contents

The Demographic Tide: A Complete Guide to Understanding How Retirement Flows Will Reshape Equity Markets ..... 1

A Teaching Companion for the BDBA Capstone Project ..... 1

Table of Contents ..... 1

Chapter 1: How Financial Markets Work ..... 4

What Is a Financial Market? ..... 4

Prices, Supply, and Demand — The Engine of Everything ..... 4

What Is a Stock? ..... 5

What Determines a Stock Price? ..... 6

View 1: The Fundamental Value Story ..... 6

View 2: The Supply-and-Demand Story ..... 7

Stocks vs. Bonds ..... 7

What Is "The Stock Market"? ..... 8

How Trading Works: Primary vs. Secondary Markets ..... 9

Market Capitalization and Why It Matters ..... 10

What "Elasticity" Means (A Preview) ..... 10

Chapter 1 Summary ..... 11
Chapter 2: The Life-Cycle Savings Story ..... 12
Why Do People Save? ..... 12
Human Capital vs. Financial Capital ..... 13
The Overlapping Generations Framework ..... 13
The Baby Boom: A Demographic Earthquake ..... 14
Bequests: Do They Break the Story? ..... 15
The Favored Cohort Effect ..... 17
Chapter 2 Summary ..... 17
Chapter 3: How American Retirement Systems Work ..... 18
Why the Institutional Details Matter ..... 18
A Brief History: From Pensions to 401(k)s ..... 19
How a 401(k) Actually Works ..... 20
Auto-Enrollment: The Behavioral Revolution ..... 20
Target-Date Funds: The Glide Path Machine ..... 22
Required Minimum Distributions (RMDs) ..... 24
Voluntary Withdrawals and Actual Drawdown Behavior ..... 25
SECURE 2.0 and the Regulatory Landscape ..... 26
How All the Pieces Fit Together: The Inflow-to-Outflow Pipeline ..... 27
Chapter 3 Summary ..... 28
Chapter 4: Do Demographics Move Markets? The Great Debate ..... 29
The Question That Haunts This Capstone ..... 29
Poterba (2001): “Demographic Structure and Asset Returns” — The Null Result ..... 29
Geanakoplos, Magill & Quinzii (2004): The MY Ratio ..... 31
The Brooks Critique: The Strongest Objection ..... 32
Abel (2001) Revisited: General Equilibrium Discipline ..... 34
Poterba (2001) and Geanakoplos (2004) — The Synthesis ..... 35
Chapter 4 Summary ..... 35

Chapter 5: The Inelastic Markets Revolution ..... 36
Why This Chapter Is the Heart of Everything ..... 36
What Does "Inelastic" Mean in This Context? ..... 37
The Undergraduate Example: A Two-Fund World ..... 38
Where Does Elasticity $\zeta = 0.2$ Come From? ..... 40
Flows Explain Most of Stock Market Fluctuations ..... 42
The General Equilibrium Model: Equity Premium, Predictability, and the Pricing Kernel ..... 43
Corporate Buybacks: The Key Counterforce ..... 44
Government "Equity QE" — A Policy Digression ..... 45
What About Slow-Moving Capital? ..... 46
ETFs and Volatility Amplification ..... 46
The True Passive Share: Chinco and Sammon (2024) ..... 47
Beliefs and Portfolios: Giglio et al. (2021) ..... 48
Chapter 5 Summary ..... 49
Chapter 6: How Competitive Is the Stock Market Really? ..... 49
Why We Need a Second Paper ..... 49
The Demand Systems Approach ..... 50
The Strategic Substitutability Result: $\chi = 3$ ..... 50
The Pass-Through Formula ..... 51
Russell Reconstitution: A Natural Experiment ..... 52
Two Methods, Same Conclusion ..... 53
The Time-Varying Multiplier: How We Use Haddad et al. in Module 5 ..... 53
Chapter 6 Summary ..... 54
Chapter 7: Synthesis — How All the Pieces Connect ..... 55
The Intellectual Arc ..... 55
Mapping Papers to Modules ..... 56
Module 1: Demographics (U.S. Population by Age Cohort, 2025-2045) ..... 56
Module 2: Inflows (Annual Equity Contributions by Age Cohort) ..... 56

Module 3: Outflows (Annual Equity Withdrawals by Age Cohort) ..... 56
Module 4: Net Flow Aggregation..... 57
Module 5: Price Impact Application ..... 57
The Cross-Paper Connections That Make This Work ..... 57
The Cascading Problems of Inelasticity ..... 58
Corporate Buybacks as Counterforce: A Detailed Treatment ..... 59
The Value Investing Paradox ..... 61
International Synchronization: Why Foreign Investors Will Not Save Us ..... 61
Big-Cap Concentration Effect ..... 62
The Literature Gap Our Capstone Fills ..... 63
Assumptions and Limitations ..... 63
The Framing Sentence..... 64
Chapter 7 Summary..... 65

---

# Chapter 1: How Financial Markets Work

## What Is a Financial Market?

A financial market is simply a place — physical or electronic — where people buy and sell financial assets. That sentence sounds trivially obvious, but almost everything in this capstone flows from taking it seriously. If you understand what "buying" and "selling" actually do to prices, and who the buyers and sellers are, and what happens when the composition of buyers and sellers changes, then you already have the core intuition for why demographic shifts matter for equity valuations.

Let us build the whole picture from the ground up.

## Prices, Supply, and Demand — The Engine of Everything

The price of anything — a loaf of bread, a barrel of oil, a share of Apple stock — is determined by the interaction of supply and demand. This is not a metaphor. It is the literal mechanism.

**Demand** means the total quantity of something that buyers are willing to purchase at each possible price. Generally, lower prices attract more buyers (or induce existing buyers to buy more), so demand curves slope downward.

**Supply** means the total quantity that sellers are willing to offer at each possible price. Generally, higher prices attract more sellers (or induce existing sellers to sell more), so supply curves slope upward.

The **equilibrium price** is the price at which the quantity demanded equals the quantity supplied. At this price, every willing buyer finds a willing seller, and the market "clears."

Here is why this matters for the capstone: if something causes a large group of sellers to enter the market simultaneously — say, tens of millions of retirees who need to sell their stock holdings to fund living expenses — and the number of willing buyers does not increase by the same amount, then the equilibrium price must fall. That is the core mechanism. Everything else in this document is about quantifying how large the selling pressure will be, how responsive buyers are, and therefore how much prices will move.

> ***Key Takeaway:*** *Prices are not determined by some abstract force. They are determined by the balance between people wanting to buy and people wanting to sell. Shift that balance, and prices move. The entire capstone is about a coming shift in that balance.*

## What Is a Stock?

A **stock** (also called a **share** or **equity**) is a claim on the ownership of a corporation. If a company has issued 1,000 shares and you own 100 of them, you own 10% of that company. Ownership entitles you to two things: a share of the company's profits (potentially paid out as **dividends**) and a vote in major corporate decisions (like electing the board of directors).

When people talk about "the stock market," they mean the collection of all publicly traded stocks — companies whose shares can be bought and sold by anyone through an exchange.

The **stock price** is the price of one share. If Apple's stock price is $180 and there are roughly 15 billion shares outstanding, then the total value the market assigns to all of Apple's equity is approximately $180 × 15 billion = $2.7 trillion. This total value is called the **market capitalization** (or **market cap**).

This concept — market capitalization — is absolutely central to the capstone. When Gabaix and Koijen say that $1 of flow into the stock market increases total market capitalization by approximately $5, they are saying that a small amount of actual buying pushes up the price of every existing share, and since there are many existing shares, the total value increase is much larger than the initial dollar. We will spend all of Chapter 5 building this intuition.

## What Determines a Stock Price?

There are two complementary ways to think about what determines a stock's price, and understanding both is essential for the capstone.

### View 1: The Fundamental Value Story

The first view says that a stock's price reflects the present value of all future cash flows the company will generate for shareholders. If you own a share of a company that will pay you dividends forever, the value of that share today is the sum of all those future dividends, discounted back to the present because a dollar tomorrow is worth less than a dollar today.

This is the **discounted cash flow (DCF)** approach. The key equation is:

$$Price = D_1/(1+r) + D_2/(1+r)^2 + D_3/(1+r)^3 + \ldots$$

Where $D_1$, $D_2$, $D_3$ are the expected future dividends and $r$ is the **discount rate** — the return investors require to compensate them for risk and for waiting.

Think of it this way. Suppose a company will pay $1 in dividends every year forever, and investors require a 5% annual return. Then:

$$Price = \$1/0.05 = \$20$$

This is a **perpetuity formula** — when dividends are constant forever, the price is just the dividend divided by the discount rate.

Notice something crucial: the price goes up if expected future dividends go up (numerator gets bigger), or if the discount rate goes down (denominator gets smaller). The discount rate is often called the **required return** or the **cost of equity capital**.

When investors demand less compensation for holding stocks — perhaps because they perceive less risk — stock prices rise even if nothing about the company's actual business has changed.

### View 2: The Supply-and-Demand Story

The second view says that a stock's price at any given moment is simply the price at which the last trade occurred — the price at which a buyer and seller agreed. This is a **market microstructure** view.

These two views are not contradictory. In a perfectly efficient market with infinitely many sophisticated investors, the supply-and-demand price converges to the fundamental value. Every investor would calculate the DCF value, and the market price would equal it.

But here is the critical insight for the capstone: **markets are not perfectly efficient, and investors are not infinitely elastic.** Many of the largest participants in today's stock market — passive index funds, target-date funds, 401(k) plans with auto-enrollment — do not calculate fundamental values at all. They simply buy whatever the index holds, in proportion to its weight, regardless of price. When these price-insensitive investors dominate the market, the supply-and-demand mechanism operates differently than textbooks assume. Flows of money into and out of the market have outsized effects on prices, because there are not enough price-sensitive investors to absorb the imbalance.

This is exactly the world Gabaix and Koijen (2024) describe, and it is the world in which demographic outflows will operate.

> ***Key Takeaway:*** *Stock prices reflect both fundamental value and the balance of buying and selling pressure. When price-insensitive investors dominate, the second force becomes relatively more powerful. This is the environment in which demographic retirement flows will play out.*
>
> ***Examiner Alert:*** *An examiner could challenge the claim that markets have shifted from efficient pricing to flow-driven pricing. The defense is that this is not an either/or — both forces always operate. The question is the relative weight, and Gabaix-Koijen and Haddad et al. provide independent empirical evidence that the flow channel has become economically large. We address this in Chapters 5 and 6.*

## Stocks vs. Bonds

A **bond** is a loan. When you buy a bond, you are lending money to the issuer (a government or corporation) in exchange for regular interest payments (the **coupon**) and the return of your principal at a specified date (the **maturity date**). A U.S. Treasury bond is considered virtually risk-free because the U.S. government can always print dollars to repay its debts.

Stocks and bonds are the two primary asset classes. The key difference: bonds give you a fixed, predictable stream of cash flows (unless the issuer defaults), while stocks give you a variable, uncertain stream that depends on how well the company performs. Because stocks are riskier, investors demand higher expected returns from stocks — the difference between the expected return on stocks and the return on bonds is called the **equity risk premium**.

For the capstone, the stock-bond distinction matters because retirement savings are split between the two. A typical target-date fund (TDF) for a 25-year-old might hold 90% stocks and 10% bonds. For a 65-year-old near retirement, the same fund might hold 40% stocks and 60% bonds. This gradual shift from stocks to bonds — the **glide path** — is a systematic, mechanical force that moves money out of equities as people age. Chapter 3 explains this in detail.

## What Is "The Stock Market"?

When we say "the stock market," we usually mean an aggregate measure of all (or most) publicly traded U.S. companies. The most commonly referenced measures are:

- **S&P 500**: The 500 largest U.S. companies by market capitalization. Covers roughly 80% of total U.S. stock market value. This is the index that most passive funds track.
- **Total Stock Market** (e.g., CRSP U.S. Total Market Index or Wilshire 5000): All publicly traded U.S. stocks — roughly 3,500-4,000 companies. Broader than the S&P 500.
- **Dow Jones Industrial Average**: 30 large companies. Older, less representative, but still widely cited.

For the capstone, "the stock market" effectively means the S&P 500 and total market indices, because these are what passive index funds track. When passive funds receive inflows, they buy stocks in proportion to these indices. When they experience outflows (retirees withdrawing), they sell in proportion to these indices. The selling pressure is

not random — it follows the index weights. This is why Haddad et al.'s finding about large-cap stocks having the lowest demand elasticity is directly relevant: index-heavy large caps will face the most concentrated selling pressure.

## How Trading Works: Primary vs. Secondary Markets

There are two distinct arenas in which financial assets change hands, and confusing them is a common source of error.

The **primary market** is where securities are created and sold for the first time. When a company does an Initial Public Offering (IPO), it sells newly created shares to investors, and the company receives the proceeds. The company uses this money to invest in its business — build factories, hire workers, develop products. Primary market transactions create new securities and direct money to productive enterprise.

The **secondary market** is where existing securities are traded between investors. When you buy Apple stock on the Nasdaq, you are buying it from another investor, not from Apple. Apple gets nothing from this transaction. The secondary market is like a used car market — it provides liquidity (the ability to buy and sell easily) and price discovery (a constantly updated assessment of value), but it does not directly fund the company.

Why does this distinction matter for the capstone? Because the flows we model — retirement account contributions and withdrawals — operate almost entirely through the secondary market. When a 30-year-old's 401(k) contribution buys shares of an S&P 500 index fund, the fund buys existing shares on the secondary market, pushing prices up. When a 75-year-old's Required Minimum Distribution (RMD) forces a sale, the fund sells existing shares, pushing prices down. No new stock is being created or destroyed — only the price is changing. This is the flow-price mechanism that Gabaix and Koijen formalize.

> ***Examiner Alert:*** *An examiner might argue that secondary market price changes "do not matter" because they do not directly affect corporate investment. The response is that secondary market prices matter enormously: they determine the cost of capital for firms (through the primary market price of new equity issuance), they affect household wealth and consumption, they drive pension funding ratios, and they influence corporate decisions through management compensation tied to stock prices. A sustained decline in equity valuations has real economic consequences even if no company directly "loses*

*money" from the stock price falling. This is the cascading-problems-of-inelasticity argument we develop in Chapter 7.*

## Market Capitalization and Why It Matters

Let us revisit market capitalization with a concrete example, because this concept is the bridge to the Gabaix-Koijen multiplier.

Suppose the entire U.S. stock market consists of just one company with 1,000 shares outstanding. The current price is $100 per share. Total market capitalization is $100 × 1,000 = $100,000.

Now suppose a new investor enters the market and wants to buy 10 shares. But nobody is eager to sell at $100. The buyer has to bid the price up — perhaps to $102 — before sellers start to offer their shares. A trade happens at $102.

What just happened to market capitalization? The buyer spent $102 × 10 = $1,020 to acquire 10 shares. But now the "price" of every share is $102 (the last trade price). So total market capitalization is now $102 × 1,000 = $102,000 — an increase of $2,000 from the buyer's investment of just $1,020. The buyer's inflow of ~$1,000 caused a market cap increase of ~$2,000 — roughly a 2× multiplier.

This is the multiplier effect at work. It occurs because the marginal trade sets the price for all existing shares, not just the ones being traded. In a market with many price-insensitive holders (who never sell regardless of price), the price has to move a lot to find willing sellers, so the multiplier is large. In a market with many price-sensitive traders (who happily sell when price rises even a little), the price barely has to move, so the multiplier is small.

Gabaix and Koijen find that in real U.S. equity markets, the multiplier is approximately 5× — meaning $1 of net inflow increases total market capitalization by about $5. We will unpack exactly why in Chapter 5.

> ***Key Takeaway:*** *Market capitalization changes are driven by marginal trades, not by the total stock of money in the market. A relatively small flow can move the total value of the market by multiples of that flow. This is why demographic flows, even if they are small relative to total market cap, can have large price effects.*

## What "Elasticity" Means (A Preview)

We will spend Chapters 5 and 6 on elasticity in financial markets, but the concept deserves a brief introduction now because it connects everything.

In economics, **elasticity** measures how sensitive one thing is to changes in another. The most common version is price elasticity of demand: if the price of a good rises by 1%, how much does the quantity demanded fall?

If quantity demanded drops a lot (say, 2%) for a 1% price increase, demand is called **elastic** — buyers are very sensitive to price changes. If quantity demanded barely budges (say, 0.1%) for a 1% price increase, demand is called **inelastic** — buyers are insensitive to price.

Now apply this to stocks. In a market with **elastic demand**, when a large seller (say, a retiree) dumps shares, many buyers rush in at slightly lower prices because they see a bargain. The price drops only a tiny amount before equilibrium is restored. The seller's flow has minimal price impact.

In a market with **inelastic demand**, when that same seller dumps shares, very few buyers step in. The price has to fall a lot before enough willing buyers appear. The seller's flow has a large price impact.

This is the entire ballgame. The central question of the capstone is: how inelastic is the U.S. stock market, and what does that mean for the price impact of demographic outflows?

Gabaix and Koijen (2024) answer: the market is very inelastic, with a demand elasticity of about 0.2. Haddad, Huebner, and Loualiche (2025) show that it is becoming more inelastic over time as passive investing grows. Both findings point to the same conclusion — demographic selling pressure will have amplified effects on equity prices.

> ***Connection to Our Model:*** *Elasticity is what connects Module 4 (net flows) to Module 5 (price impact). The net flow number by itself is just an accounting result — it tells you how many dollars are entering or leaving the market. Elasticity is the "translation device" that converts dollars of flow into dollars of price change. Low elasticity means the translation factor is large (the multiplier is high), and small flows produce big price movements.*

## Chapter 1 Summary

Here is what to carry forward from this chapter:

**Core concepts:** Prices are set by supply and demand in the secondary market. Marginal trades set prices for all shares, creating a multiplier effect. Price-insensitive investors (passive funds, TDFs) change the dynamics because they don't respond to price signals.

**Key terms:** Market capitalization, equity risk premium, primary vs. secondary market, elasticity, price-insensitive vs. price-sensitive investors, discounted cash flow.

**Why it matters for the capstone:** Demographic retirement flows operate through the secondary market, where the balance of buyers and sellers directly determines prices. The multiplier effect means that even modest net outflows can produce large market capitalization changes. The degree of inelasticity determines how large those changes will be.

**What comes next:** Chapter 2 explains *why* demographics create buying and selling patterns — the life-cycle savings story that generates inflows when people are young and outflows when they retire.

---

# Chapter 2: The Life-Cycle Savings Story

## Why Do People Save?

The answer seems obvious — people save so they have money later. But the *pattern* of when people save and when they spend is the engine that drives everything in this capstone.

The basic idea, formalized in economics by Franco Modigliani in the 1950s (the **life-cycle hypothesis**), is that people want to maintain a roughly stable standard of living throughout their lives, even though their income varies dramatically over time. When you are young and just starting your career, your income is low. In your prime working years (roughly 35 to 60), your income peaks. In retirement, your labor income drops to zero.

If you consumed exactly your income in every period, you would be poor when young, rich in middle age, and poor again in retirement. The life-cycle hypothesis says this is not what people want — they prefer a smoother consumption path. To achieve this, they:

- **Borrow** (or save nothing) when young and income is low
- **Save heavily** during peak earning years
- **Dissave** (spend down accumulated savings) in retirement

This pattern creates the demographic pressure on asset prices that Geanakoplos, Magill, and Quinzii (2004) formalize. When a large cohort (like the baby boomers) is in its peak saving years, there is a flood of money flowing into financial markets — driving prices up. When that same large cohort enters retirement and begins spending down its savings, money flows out of markets — driving prices down.

> ***Key Takeaway:*** *The life-cycle hypothesis is not just an abstract theory. It is the mechanism that translates population age structure into financial market flows. Every dollar of retirement saving eventually becomes a dollar of retirement spending. The question is timing, scale, and price impact.*

## Human Capital vs. Financial Capital

To understand the life-cycle pattern more deeply, think about two types of wealth.

**Human capital** is the present value of all future labor income you will earn. When you are 25, your human capital is enormous — you have 40+ years of earning ahead. When you are 65, your human capital is approximately zero (assuming you retire).

**Financial capital** is your accumulated savings — stocks, bonds, retirement accounts, real estate. When you are 25, your financial capital is close to zero. When you are 65, it is (hopefully) at its peak.

Over a lifetime, you are converting human capital into financial capital. Every paycheck, you take a portion of your labor income (human capital being "spent") and put it into your retirement account (building financial capital). Then in retirement, you reverse the process — you spend down financial capital because human capital is gone.

This conversion process is what generates the flows our capstone models. Module 2 (Inflows) captures the human-to-financial capital conversion — contributions to 401(k) plans, invested through TDFs into equities. Module 3 (Outflows) captures the financial capital drawdown — withdrawals, RMDs, and TDF de-risking that move money out of equities.

## The Overlapping Generations Framework

The life-cycle hypothesis becomes powerful for market analysis when you realize that at any given moment, multiple generations are alive simultaneously, each at different stages of the save-spend cycle. This is the **overlapping generations (OLG)** framework.

Picture the U.S. economy right now. There are roughly: - 25-year-olds just entering the workforce (low saving, net borrowers) - 45-year-olds in peak earning years (heavy savers, buying equities) - 70-year-olds in retirement (drawing down savings, selling equities)

The net flow into or out of the stock market at any moment depends on the balance across all these cohorts. If the 45-year-old cohort is unusually large (as the baby boomers were in the 1990s), the saving phase dominates and net flows are strongly positive — money floods into equities, driving prices up. If the 70-year-old cohort is unusually large (as the baby boomers are becoming now), the dissaving phase dominates and net flows turn negative.

This is the framework Geanakoplos, Magill, and Quinzii (2004) build their model on. Their paper is the subject of Chapter 4, but the basic insight is simple enough to state here: **equity prices depend on the ratio of middle-aged savers to young borrowers** — what they call the MY ratio (Middle-aged to Young).

When the MY ratio is high (lots of savers relative to borrowers), there is strong demand for equities and prices are high. When the MY ratio is low, demand weakens and prices fall. The baby boom creates a massive wave that pushes the MY ratio up (1990s boom) and then down (2020s–2030s decline).

> ***Examiner Alert:*** *An examiner might argue that the OLG framework is too stylized — real people don't save and spend in such neat patterns. The defense is threefold. First, the aggregate patterns are remarkably consistent even if individual behavior varies (the law of large numbers smooths out idiosyncratic differences). Second, the institutional mechanisms we study (auto-enrollment, TDF glide paths, RMDs) impose life-cycle patterns mechanically, regardless of individual choice. Third, Geanakoplos et al. validate the framework empirically — the MY ratio tracks U.S. equity prices surprisingly well from 1945 to 2000.*

## The Baby Boom: A Demographic Earthquake

To understand why this capstone matters right now, you need to understand the baby boom.

Between 1946 and 1964, the United States experienced an extraordinary surge in births — roughly 76 million babies were born during this 18-year period, compared to about 44 million in the preceding 18 years. This was not a gradual increase. It was a sharp demographic spike followed by a sharp decline (the "baby bust" of the 1970s, when the birth rate fell below 2.0 children per woman).

This spike created a bulge in the age distribution that has been moving through the American economy like a pig through a python. At every life stage, the boomers' sheer numbers overwhelmed the systems designed for smaller cohorts:

- In the 1950s–60s: school overcrowding, suburban sprawl
- In the 1970s–80s: labor market crowding, housing demand surge
- In the 1980s–2000s: peak saving years, massive inflows into equity markets via 401(k) plans, mutual funds, and eventually index funds
- **Now and the next 20 years: retirement, drawdown, outflows**

The baby boomers are currently between 61 and 79 years old (in 2026). The oldest boomers turned 73 in 2019 and began facing Required Minimum Distributions from their retirement accounts. The youngest boomers will not reach RMD age until 2039 (under current SECURE 2.0 rules, which set the RMD start at age 75 for those born after 1960). This means we are at the very beginning of a two-decade wave of mandatory selling.

> ***Connection to Our Model:*** *The baby boom is what makes the net flow question answerable and interesting. If every generation were the same size, inflows and outflows would roughly balance and there would be no story to tell. But the baby boom created an asymmetry — a massive cohort that generated outsized inflows during its saving years and will generate outsized outflows during its retirement. Module 1 (Demographics) captures this by projecting population by single-year age cohorts through 2045.*

## Bequests: Do They Break the Story?

One important objection to the "retirees sell stocks" narrative is that many retirees do not spend all their wealth before they die. They leave **bequests** — inheritances — to their children and grandchildren. If retirees hold onto their stocks until death and pass them to younger heirs who continue holding them, then retirement might not generate the selling pressure we predict.

This is exactly the challenge Andrew Abel addresses in his 2001 paper, "Will Bequests Attenuate the Predicted Meltdown in Stock Prices When Baby Boomers Retire?" His answer is surprising and counterintuitive: **no, bequests do not save the market from demographic pressure.**

Here is why. Think about what happens in **general equilibrium** — considering both supply and demand for capital simultaneously. When boomers are thrifty and leave large bequests, they increase the *demand* for capital (more wealth chasing assets). But they also increase the *supply* of capital (more saving funds more investment in productive capacity). Abel shows that in his model — a two-period OLG framework with Cobb-Douglas production and convex capital adjustment costs — the supply shift exactly offsets the demand shift. The price of capital (the key variable for stock market valuations) is completely invariant to the strength of the bequest motive.

In plain English: even if every baby boomer left 100% of their wealth to their kids, the predicted price pattern — a boom during the boomers' saving years followed by a decline as they retire — would be exactly the same. Bequests increase the level of capital in the economy, but they do not prevent the demographic price cycle.

Abel formalizes this with a supply-demand diagram (his Figure 1). Without bequests, equilibrium is at point E. With bequests, the demand curve shifts right to $KD_2$ (because heirs have more wealth to invest), but the supply curve also shifts right to $KS_2$ (because more saving funded more capital formation). The new equilibrium F has the same price as E but a higher capital stock. The price of capital is unchanged.

> ***Key Takeaway:*** *Bequests increase the quantity of capital but do not change its price. The demographic price cycle — up during the boom saving phase, down during the retirement drawdown phase — is robust to bequest behavior. This removes a major potential objection to the capstone thesis.*
>
> ***Examiner Alert:*** *Abel's result depends on his specific model assumptions: Cobb-Douglas production, convex adjustment costs, and a particular OLG structure. An examiner could argue that different production functions or adjustment cost structures might yield different results. The defense is that Abel's qualitative point is robust — supply-side responses to bequests offset demand-side effects — even if the exact quantitative invariance depends on functional form assumptions. For the capstone, the key implication is that bequests do not eliminate demographic selling pressure; at most they might modulate its magnitude.*

*Simplified here — we skipped the mathematical derivation of the AR(1) process for the price of capital and the formal proof that gamma (bequest parameter) drops out. If you want, we can go deeper into the heuristic vs. formal model, the convex adjustment cost function, or the Diamond economy extension.*

## The Favored Cohort Effect

Geanakoplos et al. (2004) introduce a beautiful concept called the **favored cohort effect**. It works like this:

A **large cohort** (like the baby boomers) is systematically disadvantaged in asset markets over its lifetime. When boomers entered their peak saving years in the 1980s–90s, their massive collective demand for equities drove prices up. They bought expensive. Now, as boomers enter retirement and need to sell, their massive collective supply will drive prices down. They will sell cheap.

A **small cohort** (like Generation X) is systematically advantaged. When Gen X was saving, boomers were also saving — prices were high, which was bad for Gen X's purchases. But Gen X will be selling during a period when an even larger cohort (millennials) is saving, potentially supporting prices.

The favored cohort effect is a pure **terms-of-trade** phenomenon. It has nothing to do with skill, intelligence, or investment strategy. It is a mechanical consequence of cohort size interacting with life-cycle saving. Large cohorts face unfavorable prices at every stage because their own demand moves the market against them.

For the capstone, the favored cohort effect provides important context. The baby boomers' retirement is not just a neutral event — it is the unwinding of a price bubble that the boomers themselves created through their saving behavior. The flow reversal is not random; it is the predictable second half of a demographic cycle.

***Connection to Our Model:*** *The favored cohort effect explains why the inflow-to-outflow reversal matters beyond simple flow accounting. It means prices were elevated partly because of demographic demand, and the withdrawal of that demand means prices should revert. This is the price-level channel that Poterba (2001) finds in his empirical work — demographics relate to price-dividend ratios (levels) even when they do not predict returns.*

## Chapter 2 Summary

**Core concepts:** People save when working and spend when retired (life-cycle hypothesis). Multiple generations alive simultaneously create competing flows (OLG framework). Cohort size determines which phase dominates at any point in time.

**Key terms:** Life-cycle hypothesis, human capital vs. financial capital, overlapping generations (OLG), MY ratio, baby boom, bequest motive, favored cohort effect.

**Papers introduced:** Geanakoplos, Magill & Quinzii (2004) — MY ratio and favored cohort; Abel (2001) — bequests don't attenuate the meltdown.

**Why it matters for the capstone:** The life-cycle saving pattern, applied to the baby boom generation, predicts a reversal from net equity inflows to net equity outflows over the next two decades. This is the demographic force our model quantifies.

**What comes next:** Chapter 3 explains the *institutional machinery* through which life-cycle saving actually operates in America — 401(k) plans, auto-enrollment, target-date funds, and Required Minimum Distributions.

---

# Chapter 3: How American Retirement Systems Work

## Why the Institutional Details Matter

Chapter 2 told the theoretical story: people save when working and spend when retired, and cohort size determines which phase dominates. But theory alone does not tell you the specific dollar amounts flowing in and out, or through which channels. For that, you need to understand the plumbing — the actual institutional mechanisms through which American retirement saving operates.

This matters enormously for the capstone because the flows we model are not abstract "life-cycle saving" — they are concrete, measurable, mechanically determined financial transactions. A 30-year-old's 401(k) contribution goes through a specific pathway: payroll deduction → plan custodian → target-date fund → equity index fund → secondary market stock purchase. A 75-year-old's Required Minimum Distribution follows an equally specific pathway in reverse.

Understanding these pathways is what makes the capstone an *accounting model* rather than a theoretical exercise. Every parameter in Modules 2 and 3 corresponds to a real institutional feature with real data behind it.

## A Brief History: From Pensions to 401(k)s

For most of the twentieth century, American workers who had employer-sponsored retirement plans had **defined benefit (DB) pensions**. Under a DB plan, the employer promises a specific monthly payment in retirement (for example, 1.5% of your final salary for each year of service). The employer, not the employee, bears the investment risk. The employer hires professional asset managers who decide how to invest the pension fund's assets. Employees make no investment decisions and bear no direct market risk.

Starting in the late 1970s and accelerating through the 1980s–2000s, American retirement shifted dramatically toward **defined contribution (DC) plans**, primarily the **401(k)**. Under a DC plan, both the employer and employee contribute to an individual account owned by the employee. The employee chooses how to invest these contributions (from a menu of options offered by the plan). The employee bears the investment risk — if the stock market goes up, their retirement account grows; if it goes down, their account shrinks.

This shift matters for the capstone for three reasons:

**First**, the DB-to-DC shift moved investment decisions from professional asset managers to individual employees. Professional managers respond to valuations — they are more likely to be price-sensitive. Individual employees, as we will see, are overwhelmingly passive and inertial — they are price-insensitive. The shift expanded the pool of price-insensitive market participants, which is a key driver of the inelasticity story in Chapters 5 and 6.

**Second**, DC plans create visible, trackable flows that can be modeled. Every dollar contributed to a 401(k) is a measurable inflow. Every dollar withdrawn or distributed is a measurable outflow. DB plans are murkier because the employer's pension fund makes aggregate investment decisions that are harder to map to individual cohort behavior.

**Third**, DC plan assets are held in individual accounts that must eventually be distributed to the owner (or their heirs). Unlike DB plans, where the employer can

reinvest assets across generations indefinitely, DC assets have a finite lifespan — they must come out. This is where RMDs enter the picture.

## How a 401(k) Actually Works

Here is the step-by-step mechanics of how money flows through a 401(k) plan. Every step corresponds to a parameter in our capstone model.

**Step 1: Eligibility and enrollment.** An employer offers a 401(k) plan. Employees meeting certain criteria (typically age 21+ and 1 year of service) become eligible. Before 2006, employees had to actively enroll — they had to fill out forms, choose a contribution rate, and select investment options. Many never bothered. Participation rates at many companies hovered around 40-60%.

**Step 2: Contribution.** Once enrolled, the employee designates a percentage of each paycheck to contribute to the 401(k). The employer may match a portion (for example, 50% match up to 6% of salary). Contributions are typically pre-tax, meaning they reduce current taxable income. The IRS sets annual limits ($23,500 in 2025, plus $7,500 catch-up for those 50+).

**Step 3: Investment allocation.** The contributed money must be invested in one or more options from the plan's menu. Typical menus include stock index funds, bond index funds, money market funds, company stock, and target-date funds. The employee's choice here determines what fraction of their contribution ends up in equities — which is what Module 2 (Inflows) needs to know.

**Step 4: Accumulation.** Over decades of working, contributions compound with investment returns. The employee's account grows. During this phase, the account is generating demand for equities (inflow).

**Step 5: Distribution.** Upon retirement (or after reaching certain ages), the employee begins withdrawing money from the account. Each withdrawal requires selling assets — if the account holds stocks, stocks must be sold. This generates supply of equities (outflow). If the employee does not withdraw voluntarily, the IRS eventually forces withdrawals through RMDs.

## Auto-Enrollment: The Behavioral Revolution

The single most consequential change in American retirement saving in the past quarter-century was the adoption of **automatic enrollment**. Understanding it is essential because it transformed the flow pipeline our model tracks.

Before auto-enrollment, 401(k) plans required employees to actively opt *in*. They had to fill out paperwork, choose a contribution rate, and select investments. Behavioral economics predicts — and Madrian and Shea (2001) prove empirically — that many people never get around to it, even when it is strongly in their financial interest.

Madrian and Shea study a large U.S. corporation (anonymized in the paper as a health-care products company with roughly 30,000+ employees) that switched from opt-in to automatic enrollment on April 1, 1998. Their findings are striking:

- **Participation jumped from roughly 37% to 86%** among newly eligible employees. This is not a small effect — it is a doubling.
- **71% of auto-enrolled employees stuck with the default contribution rate** (3% of salary) even after more than three years, even though the company match incentivized contributing at least 6%.
- **The vast majority stuck with the default investment fund** — at the time, a money market fund (extremely conservative). The default was not objectively optimal, but people accepted it because it required no action.

The mechanism is **inertia** — the tendency to stick with whatever is already in place. People don't actively choose to save 3% in the default fund; they simply never get around to changing it. Madrian and Shea show this is not because employees are satisfied with the default. When surveyed, many said they intended to change their allocation but hadn't yet. The "power of suggestion" in the paper's title refers to the finding that employees interpret the employer's default as implicit advice — "if the company set it to 3%, that's probably what I should do."

The policy implication was enormous. The **Pension Protection Act of 2006 (PPA)** cleared the legal obstacles for employers to auto-enroll workers and designated target-date funds (TDFs) as a **qualified default investment alternative (QDIA)**. This meant employers could auto-enroll workers and put their contributions into TDFs without fear of fiduciary liability.

The result was an institutional pipeline that operates almost entirely without active decision-making:

1. Worker is hired → automatically enrolled in 401(k)

2. Contribution rate set to default (typically 3-6%) → worker does not change it
3. Contributions directed to default TDF based on birth year → worker does not change it
4. TDF invests according to mechanical glide path → no human decision involved

Every step of this pipeline creates **price-insensitive flow**. The money goes into equities not because anyone evaluated stock valuations and decided stocks were attractively priced, but because the institutional machinery directed it there by default. This is exactly the type of flow that Gabaix and Koijen (2024) argue makes markets inelastic.

> ***Connection to Our Model:*** *Madrian and Shea provide the behavioral foundation for Module 2 (Inflows). The ~85% participation rate under auto-enrollment (now standard across most large and medium employers, and mandated for new plans under SECURE 2.0) tells us what fraction of eligible workers contribute. The default contribution rate tells us how much. The default allocation to TDFs tells us where the money goes. Combined with Census population data by age, these parameters let us calculate annual equity inflows by cohort.*
>
> ***Examiner Alert:*** *An examiner might argue that auto-enrollment's effects could diminish over time — as employees become more financially literate, they might override defaults more often. The defense is that the empirical evidence shows the opposite. Madrian and Shea's finding of persistent default behavior (71% still at default after 3+ years) has been replicated in numerous subsequent studies. Giglio et al. (2021) provide further evidence: even when investors change their beliefs about markets dramatically, their portfolios barely move. Inertia is not a temporary phenomenon — it is a durable feature of human behavior.*

## Target-Date Funds: The Glide Path Machine

**Target-date funds (TDFs)** are now the dominant vehicle for retirement saving in America, and understanding how they work is critical for the capstone's flow calculations.

A TDF is a fund-of-funds that maintains a specific allocation between stocks and bonds, adjusting that allocation automatically over time according to a predetermined **glide path** tied to the investor's expected retirement date. You buy the fund named

with the year you plan to retire (for example, a 25-year-old in 2025 might buy "Target Date 2060"), and the fund does everything else.

Parker, Schoar, and Sun (2020) provide the most comprehensive study of TDF market dynamics. Key facts from their paper:

**TDF growth has been explosive.** TDF and balanced fund assets rose from under $8 billion in 2000 to almost $6 trillion by 2021, representing about 22% of the $27 trillion held in U.S. mutual funds. Of this, about $3.2 trillion was in TDFs specifically, including $1.4 trillion in target-date collective investment trusts (CITs) used by large employers.

**Glide paths follow a standard pattern.** An average TDF allocates 80-90% to equity funds and 10-20% to bond funds until about 25 years before the target date. Then the equity share declines smoothly, reaching 30-40% about 10 years after the target date. This means:

- A 25-year-old in a 2065 TDF: ~90% equities
- A 40-year-old in a 2050 TDF: ~85% equities
- A 55-year-old in a 2035 TDF: ~60% equities
- A 65-year-old at the target date: ~45% equities
- A 75-year-old, 10 years past target: ~30% equities

**TDF rebalancing is mechanical and macrocontrarian.** This is one of Parker et al.'s most important findings. When the stock market rises, TDFs must sell stocks to return to their target allocation. When the stock market falls, they must buy stocks. This rebalancing happens quickly — roughly 45% within the same month, 25% the following month, and 10% two months later. Passive TDFs (holding index funds) rebalance especially fast, sometimes within a day.

This macrocontrarian behavior means TDFs act as automatic stabilizers for their target equity allocation. They do not respond to valuations, sentiment, or economic conditions — they respond only to deviations from their prescribed allocation. This is pure price-insensitive behavior, contributing directly to the inelasticity of markets.

**TDFs create a systematic equity-to-bond flow as participants age.** As each TDF approaches and passes its target date, its equity allocation automatically declines. This is not a one-time event — it is a continuous process. Every year, the glide path moves slightly toward bonds, and the fund mechanically sells equities and buys bonds.

For our capstone, this is a separate outflow channel (TDF de-risking) in addition to RMDs and voluntary withdrawals.

> ***Key Takeaway:*** *TDFs transform the retirement savings pipeline into a fully automated system. Money enters via auto-enrolled payroll contributions, is invested according to a mechanical glide path with no human valuation judgment, and systematically shifts from equities to bonds over a 40-year horizon. This entire process operates without anyone making a price-sensitive investment decision.*

## Required Minimum Distributions (RMDs)

RMDs are the mandatory withdrawals from tax-deferred retirement accounts (traditional 401(k)s and IRAs) that the IRS requires once account holders reach a certain age. They are the most mechanically predictable outflow channel in our model because they are legally mandated, formulaically determined, and impossible to avoid (short of converting the entire account to a Roth IRA, which triggers a massive upfront tax bill).

**How RMDs work:** The IRS Uniform Lifetime Table specifies a "distribution period" (a divisor) for each age. To calculate the RMD, you divide the account balance on December 31 of the prior year by the distribution period. The result is the minimum amount that must be withdrawn during the year.

For example, at age 75, the distribution period is approximately 24.6. If a retiree has a $500,000 IRA, the RMD is $500,000 ÷ 24.6 ≈ $20,325. At age 80, the distribution period drops to about 20.2, so the same balance would require $500,000 ÷ 20.2 ≈ $24,752. The distribution period decreases every year, meaning the percentage withdrawn increases with age — roughly 3.7% at 73, 4.1% at 75, 5.0% at 80, 6.3% at 85, and 8.8% at 90.

**SECURE Act and SECURE 2.0:** The RMD starting age has been pushed back by recent legislation: - Before 2020: age 70½ - SECURE Act (2020): age 72 (for those born 1950-1950) - SECURE 2.0 (2022): age 73 (born 1951–1959) and age 75 (born 1960 and later)

For the capstone, the key implication is that the youngest baby boomers (born 1964) will not face RMDs until 2039, and the peak of boomer RMD withdrawals will extend well into the late 2030s and early 2040s. This is a longer timeline than casual discussion suggests — the outflow wave does not end with the oldest boomers.

**RMDs as forced selling:** Here is the critical mechanism. Most retirees hold their retirement accounts in mutual funds or index funds. When an RMD is taken, the custodian must sell fund shares to generate the cash distribution. If the account is invested in an S&P 500 index fund, the fund sells S&P 500 stocks in proportion to index weights. This is involuntary, mechanical equity selling — the retiree has no choice about whether to sell (they must take the RMD) and typically no choice about what to sell (the fund sells pro-rata).

> ***Connection to Our Model:*** *RMDs are Module 3's most reliable parameter. For each age cohort, we know: (a) how many people are in that cohort (Module 1), (b) what fraction have DC retirement accounts with balances, (c) the average account balance, and (d) the exact RMD percentage (from the IRS table). The only uncertainty is in account balances (which depend on past returns and saving behavior), not in the withdrawal rule itself.*
>
> ***Examiner Alert:*** *An examiner might argue that many retirees withdraw more than the RMD minimum (voluntary withdrawals), so RMDs are not the binding constraint. This is correct and does not undermine the model — it strengthens it. RMDs set the floor for outflows; actual outflows are higher. Poterba, Venti, and Wise document that actual drawdown rates often exceed RMDs, particularly for lower-wealth households that need the money for living expenses. Our model accounts for this by including voluntary withdrawals as a separate outflow channel calibrated against empirical withdrawal data.*

## Voluntary Withdrawals and Actual Drawdown Behavior

Not all retirement outflows come from RMDs. Many retirees begin withdrawing before the RMD age (you can withdraw from a 401(k) or IRA without penalty after age 59½). Others withdraw more than the RMD requires because they need the money for living expenses, medical costs, or housing.

Poterba, Venti, and Wise study the actual composition and drawdown of wealth in retirement using the Health and Retirement Study (HRS), a comprehensive longitudinal survey of Americans over age 50. Their data provides a crucial reality check on the capstone's outflow assumptions.

Key findings from their research on household balance sheets for retirees aged 65-69:

- **Personal retirement accounts** (401(k)s and IRAs) represent about 13.9% of total household wealth for this age group, with a mean balance of roughly $121,000
- **Social Security and DB pensions** together represent about 34.9% of total wealth
- Only about 52% of these households have any personal retirement account balance at all
- Wealth is highly concentrated: the top quintile holds dramatically more than the bottom

For the capstone, this data tells us that retirement account drawdown is not uniform — wealthier households have much larger balances and therefore generate much larger absolute flows when they withdraw. The aggregate outflow is driven disproportionately by the top 20-30% of the wealth distribution.

> *Simplified here — we focused on the key numbers relevant to the model. Poterba/Venti/Wise also analyze housing wealth, DB pension annuities, and cohort effects in wealth accumulation. If you want, we can go deeper into their full wealth decomposition or their comparison across HRS waves.*

## SECURE 2.0 and the Regulatory Landscape

The SECURE 2.0 Act of 2022 is the most recent major piece of retirement legislation, and several provisions directly affect our model's parameters:

**RMD age increase to 75:** As noted above, those born in 1960 or later will not face RMDs until age 75. This delays the onset of mandatory outflows for the youngest boomers and for all subsequent cohorts.

**Mandatory auto-enrollment for new plans:** Starting in 2025, new 401(k) and 403(b) plans must auto-enroll eligible employees at a contribution rate of at least 3% (with annual auto-escalation of 1% per year up to at least 10%). This codifies into law the behavioral finding from Madrian and Shea — defaults are now the federal standard, not just best practice.

**Roth catch-up contributions:** Starting in 2026, high earners (over $145,000) must make catch-up contributions to Roth accounts (after-tax). This is relevant because Roth accounts have no RMDs during the owner's lifetime — converting assets to Roth reduces future mandatory outflows.

**Student loan matching:** Employers can match student loan payments as if they were 401(k) contributions. This may increase effective participation among younger workers burdened by student debt.

For the capstone, the most important provisions are the RMD age change (which we incorporate directly into Module 3 timing) and the auto-enrollment mandate (which affects future inflow projections in Module 2).

## How All the Pieces Fit Together: The Inflow-to-Outflow Pipeline

Let us now trace the complete lifecycle of a dollar through the American retirement system, as it relates to our capstone model:

**Age 25 — Entry:** Maria starts her first job at a company with a 401(k) plan. She is automatically enrolled at a 4% contribution rate. Her contributions are automatically directed to a Target Date 2065 fund, which holds ~90% equities. Her annual salary is $50,000, so she contributes $2,000/year. Of this, ~$1,800 flows into equity index funds. (Module 2: Inflow)

**Age 45 — Peak saving:** Maria now earns $90,000. She has never changed her contribution rate from the auto-escalated 10%. She contributes $9,000/year. Her TDF (now "Target Date 2045") holds ~70% equities, so ~$6,300 flows to equities annually. Her account balance is ~$250,000 and growing. (Module 2: still inflow, but TDF de-risking has begun shifting the equity share down)

**Age 55 — Late career:** Maria earns $100,000 and contributes $10,000 + $7,500 catch-up = $17,500. Her TDF holds ~55% equities, so ~$9,625 flows to equities. However, the TDF itself is now *selling* equities as it de-risks — perhaps $5,000 per year as the glide path shifts from 55% to 50% equity. Net new equity exposure from Maria's account: ~$4,625. (Module 2 inflow + Module 3 TDF de-risking outflow beginning)

**Age 65 — Retirement:** Maria stops contributing. Her account is ~$800,000. The TDF holds ~40% equities (~$320,000 in equities). She begins voluntary withdrawals of ~$30,000/year. The TDF continues de-risking toward ~35% equity. Both the withdrawal and the de-risking generate equity outflows. (Module 3: outflow begins)

**Age 75 — RMDs begin:** Maria's account is ~$600,000. The RMD is ~$24,000 (4%). Her TDF holds ~30% equities. The RMD forces a sale, including ~$7,200 in equity

sales from the equity portion. She may also take voluntary withdrawals above the RMD. (Module 3: mandatory outflow, increasing annually)

**Age 85 — Late retirement:** Maria's account has drawn down to ~$300,000. The RMD is now ~$24,000 (8% rate at 85). Equity allocation in the TDF is ~25%. Annual equity outflow from her account: RMD equity portion (~$6,000) plus any voluntary excess. (Module 3: outflow accelerating as RMD percentage rises)

Now multiply Maria by millions of baby boomers, all moving through this pipeline on slightly different timelines. The aggregate flow pattern is the sum of all these individual trajectories, and it is what our model calculates.

> ***Key Takeaway:*** *The American retirement system has been engineered — through auto-enrollment, TDF defaults, and RMD rules — to create mechanically predictable, price-insensitive flows into and out of equity markets. These flows are not driven by investor sentiment or valuation judgments. They are driven by demographic age structure, institutional defaults, and tax law. This is what makes them modelable.*

## Chapter 3 Summary

**Core concepts:** The DB-to-DC shift put investment decisions in the hands of passive, inertial individuals. Auto-enrollment dramatically increased participation and cemented TDFs as the default. TDF glide paths mechanically shift assets from equities to bonds over a worker's lifetime. RMDs force mandatory withdrawals from tax-deferred accounts starting at 73 or 75.

**Key terms:** 401(k), defined benefit vs. defined contribution, auto-enrollment, default effect, target-date fund (TDF), glide path, QDIA, Required Minimum Distribution (RMD), SECURE 2.0.

**Papers drawn from:** Madrian & Shea (2001) — auto-enrollment and inertia; Parker, Schoar & Sun (2020) — TDF growth and rebalancing dynamics; Poterba, Venti & Wise — retirement wealth composition and drawdown.

**Why it matters for the capstone:** Every parameter in Modules 2 and 3 maps to a specific institutional mechanism described in this chapter. Participation rates come from auto-enrollment effects. Equity allocation comes from TDF glide paths. Outflow timing comes from RMD rules and voluntary withdrawal data. The institutional machinery makes the flows predictable and modelable.

**What comes next:** Chapter 4 asks the big question — given these demographic flows, do they actually move market prices? We examine the three-decade academic debate over whether demographics affect asset returns and prices.

---

# Chapter 4: Do Demographics Move Markets? The Great Debate

## The Question That Haunts This Capstone

We have established that baby boomers' life-cycle saving behavior generates massive flows into equity markets (during working years) and out of them (during retirement). We have described the institutional plumbing through which these flows operate. But none of that matters if demographic flows do not actually move prices.

This is not a hypothetical concern. There is a genuine academic debate — stretching back to the late 1990s — about whether demographics affect asset prices. Some prominent economists find little evidence. Others find strong evidence. The capstone cannot simply ignore the skeptics; it must engage honestly with the null results and explain why the thesis survives them.

This chapter walks through the three core papers in the demographic debate: Poterba (2001), who finds a weak empirical relationship; Abel (2001), who shows bequests don't save the market; and Geanakoplos, Magill, and Quinzii (2004), who build the most comprehensive theoretical framework linking demographics to prices. We also address the strongest critique of the entire literature — Robin Brooks' discussant comment on Geanakoplos et al.

## Poterba (2001): "Demographic Structure and Asset Returns" — The Null Result

James Poterba's paper is the most important challenge the capstone faces. Published in the *Review of Economics and Statistics*, it asks a direct empirical question: does population age structure predict asset returns?

**What Poterba does:** He constructs a "projected asset demand" variable from Survey of Consumer Finances (SCF) data, combining age-specific wealth-holding profiles with Census population projections. He then tests whether this variable predicts stock returns, bond returns, price-dividend ratios, and Treasury bill rates using postwar U.S. data.

**What he finds:** The relationship between demographics and *returns* is statistically weak. In most specifications, the age-structure variable does not significantly predict stock returns or bond returns. However — and this is crucial — Poterba does find that his projected asset demand variable relates to **price-dividend ratios** (P/D ratios). When demographics predict high asset demand (lots of middle-aged savers), P/D ratios are indeed elevated.

**Why this seemingly undermines the capstone:** If demographics don't predict returns, then maybe they don't matter for markets. An examiner could cite Poterba as evidence that the entire demographic-market-impact thesis is empirically unsupported.

**Five reasons Poterba does NOT undermine the capstone:**

**Reason 1: Returns vs. levels.** Poterba tests *returns* (percentage changes), but the Gabaix-Koijen multiplier operates on *price levels*. These are fundamentally different. A demographic force can permanently shift the price level downward (lowering P/D ratios) without generating predictable returns. Think of it this way: if everyone expects prices to fall gradually due to demographics, the market prices this in, and the actual decline generates no abnormal return — yet the price level is still lower than it would have been without demographic pressure. Poterba's own finding that demographics track P/D ratios is actually consistent with the capstone thesis.

**Reason 2: The wrong channel.** Poterba tests the *aggregate* demographic channel — slow-moving population ratios. The capstone models the *institutional* channel — specific flows through 401(k)s, TDFs, and RMDs. These institutional flows are sharper, more concentrated, and more recent (post-2006, after the PPA created the TDF pipeline). Poterba's sample period predates the institutional structure the capstone analyzes.

**Reason 3: No passive fund dominance.** Poterba's data comes from an era when active management dominated. The inelasticity argument (Gabaix-Koijen) depends on passive, price-insensitive investors dominating the market, which only became true in the 2010s. Testing demographic effects in a market dominated by active, price-

sensitive investors is like testing whether a heavy truck will break a wooden bridge by testing with a bicycle — the forces are qualitatively different.

**Reason 4: Poterba himself finds the level effect.** In his own results, projected asset demand does predict P/D ratios. He downplays this finding because his paper focuses on returns, but it is exactly the channel the capstone thesis operates through.

**Reason 5: His conclusion calls for what we're doing.** Poterba concludes by noting that the aggregate approach may miss specific channels through which demographics affect asset markets. He explicitly calls for research that models the institutional mechanisms. That is precisely what the capstone does.

> ***Examiner Alert:*** *An examiner who reads Poterba superficially might say "demographics don't predict returns, end of story." The prepared response is: "Poterba shows demographics don't predict returns through aggregate channels in a pre-passive-era sample. Our model operates on price levels through institutional channels in the current passive-dominated market structure. The operating environment has fundamentally changed since Poterba's analysis. Moreover, Poterba does find the price-level effect that aligns with our thesis."*

## Geanakoplos, Magill & Quinzii (2004): The MY Ratio

While Poterba attacks the question empirically, Geanakoplos, Magill, and Quinzii (GMQ) attack it theoretically using a rigorous overlapping generations model published in *Brookings Papers on Economic Activity*.

**The core mechanism:** GMQ build an OLG model where three age groups coexist: the Young (who borrow), the Middle-aged (who save by buying equities), and the Old (who sell equities to fund retirement). The equity price is determined by the balance between middle-aged buyers and young borrowers. GMQ define the **MY ratio** as the population of 40-49 year olds divided by the population of 20-29 year olds.

**The key prediction:** When the MY ratio is high — lots of middle-aged savers relative to young borrowers — there is strong demand for equities and P/E ratios are high. When the MY ratio falls, demand weakens and P/E ratios decline.

**Empirical validation:** Here is the remarkable finding. The MY ratio tracks the real S&P 500 price index surprisingly well from 1945 to 2000. The post-war boom, the 1970s stagnation, and the 1990s bull market all align with movements in the MY ratio.

GMQ calibrate their model with a risk aversion parameter of $\alpha = 4$ and find P/E ratios of ~19 when a large cohort is middle-aged versus ~8 when a large cohort is young.

**Forward projection:** The MY ratio peaked around 2000 (when the boomers were 36-54) and has been declining since. GMQ project continued decline through the 2020s, implying downward pressure on P/E ratios. This aligns with the capstone's timing — the demographic headwind began when the boomers started exiting their peak saving years.

**Exchange economy vs. production economy:** GMQ analyze both cases. In an exchange economy (fixed supply of equity, no capital accumulation), demographic effects on prices are maximized. In a production economy (where saving funds new capital investment), effects are smaller because some of the excess saving goes into capital formation rather than bidding up equity prices. Importantly, even in the production economy, the P/E ratio is still sensitive to demographics — the damping is modest, not complete.

> ***Key Takeaway:*** *The MY ratio provides an independent validation framework for the capstone. It predicts that equity valuations should be under pressure when a large cohort exits the peak saving age bracket — which is exactly what is happening now and will intensify through 2040.*
>
> ***Connection to Our Model:*** *The MY ratio captures the same fundamental force our model quantifies — the balance of buying and selling across age cohorts. The difference is that our model goes deeper: instead of using aggregate population ratios, we trace specific institutional flows (401(k) contributions, TDF allocations, RMDs) to produce dollar-denominated net flow projections. The MY ratio is the concept; our model is the implementation.*

## The Brooks Critique: The Strongest Objection

Robin Brooks, serving as a discussant on the Geanakoplos et al. paper, raised the single most damaging critique of the demographic approach. Every examiner will be familiar with some version of this argument, so understanding it — and having a response — is essential.

**The critique:** Brooks points out that the MY ratio's fit to equity prices over 1945-2000 may be partly coincidental because several major structural changes in the U.S. economy occurred on similar timescales as the baby boom demographic wave:

- **Social Security expansion** (1950s-70s): Expanded coverage and benefit levels changed saving behavior across the same decades the boomers were growing up
- **Financial deepening** (1980s-90s): Deregulation, the creation of money market funds, and easier access to financial markets coincided with boomer entry into peak saving
- **401(k) creation and growth** (1980s-2000s): The DC revolution happened exactly when boomers started saving heavily

The concern is that the MY ratio and these structural changes are correlated in time but driven by different causes. If the 1990s bull market was driven by financial innovation and deregulation rather than demographics, then the MY ratio's predictive power may not apply to the future.

**The response (four parts):**

**Part 1: Even if the fit is partly coincidental, the mechanism is real.** The *mechanism* — middle-aged people save more than young people, and this creates net demand for equities — is not in dispute. The only question is the magnitude. Even if structural changes account for half of the MY ratio's historical fit, demographics still account for the other half, and that is enough to generate meaningful price impacts through the Gabaix-Koijen multiplier.

**Part 2: The institutional machinery is now transparent.** Brooks' critique applies to the *indirect* channel (MY ratio as a reduced-form predictor). Our capstone models the *direct* channel — dollar flows through TDFs, 401(k)s, and RMDs. We can observe these flows in real time. The structural-change confound disappears when you track the actual money rather than relying on population ratios as a proxy.

**Part 3: The structural changes reinforced demographics, not competed with them.** The creation of 401(k) plans did not replace demographic effects — it *amplified* them. Auto-enrollment, TDFs, and passive indexing channeled demographic saving more directly into equities than the pre-1980 system did. The future structural landscape (continued passive growth, SECURE 2.0 mandates) will similarly amplify demographic outflows.

**Part 4: The inelastic markets framework makes this testable.** Brooks' critique predates Gabaix-Koijen (2024) and Haddad et al. (2025). The inelastic markets hypothesis provides a separate, independently validated mechanism (the 5x multiplier) through which flows affect prices. Even if you reject the MY ratio as a

forecasting tool, the Gabaix-Koijen multiplier gives you a direct, empirically estimated translation from flows to price impacts. The capstone does not depend on the MY ratio — it uses it for context and then applies a different, more robust methodology.

> ***Examiner Alert:*** *The Brooks critique is the number-one thing an examiner will reach for to challenge the demographic thesis. The strongest response is: "We agree the MY ratio alone is an imperfect forecasting tool due to potential confounds with structural change. That is precisely why our model does not rely on the MY ratio. Instead, we trace specific institutional flows and apply the Gabaix-Koijen multiplier — an independently estimated, causally identified price-impact coefficient. The demographic story and the inelastic markets story reinforce each other, but neither depends on the other."*

## Abel (2001) Revisited: General Equilibrium Discipline

We covered Abel's bequest result in Chapter 2, but his paper contributes something else to the debate: a reminder that partial equilibrium intuitions can be misleading.

In partial equilibrium (looking at demand alone), the story is simple: boomers sell → supply increases → price falls. But Abel shows that in general equilibrium, supply-side responses partially absorb the shock. Boomers' saving during their working years funded real capital investment. That capital now exists and generates returns. The price of capital depends on the interaction of supply and demand, not demand alone.

For the capstone, Abel's paper teaches an important lesson: **we should be cautious about treating the Gabaix-Koijen multiplier as a precise prediction for sustained outflows.** The multiplier was estimated during a period of predominantly *inflows*. Whether it applies symmetrically to sustained *outflows* over decades is an open question — in general equilibrium, price declines trigger responses (corporate buybacks, active investor rebalancing, reduced IPO supply) that partially offset the flow effect.

We return to this asymmetry question in Chapter 7 (Synthesis), where we frame it as an explicit limitation of the capstone model.

> ***Key Takeaway from Abel:*** *The capstone's flow model is inherently partial equilibrium — we project flows and apply a multiplier. General equilibrium effects (supply responses, behavioral changes, corporate actions) may dampen the price impact. This is an honest limitation, not a fatal flaw. We address it by presenting scenarios (conservative, moderate, aggressive) rather than point estimates.*

## Poterba (2001) and Geanakoplos (2004) — The Synthesis

Here is how the three papers fit together, and how the capstone navigates between them:

**Poterba** says: demographics don't predict returns (but do relate to price levels). **Translation for capstone:** We should model price-level effects (market capitalization changes), not return predictions. The Gabaix-Koijen multiplier operates on levels, not returns. ✓ Consistent.

**Abel** says: bequests don't save you, but general equilibrium dampens things. **Translation for capstone:** Don't assume the full partial-equilibrium multiplier applies without caveat. Present scenarios that bracket the uncertainty. ✓ We do this.

**Geanakoplos** says: the MY ratio predicts equity prices, but structural confounds exist. **Translation for capstone:** Use the MY ratio for contextual support and the OLG intuition for the mechanism, but don't rest the empirical identification on the MY ratio alone. Use Gabaix-Koijen for the multiplier instead. ✓ That is exactly our strategy.

The capstone sits at the intersection of all three papers, taking the mechanism from GMQ, the general equilibrium discipline from Abel, the honest engagement with null results from Poterba, and the empirical methodology from Gabaix-Koijen (Chapter 5) and Haddad et al. (Chapter 6).

## Chapter 4 Summary

**Core concepts:** The academic debate over demographics and asset prices spans three decades. Poterba finds no return predictability but does find price-level effects. GMQ build a theoretical framework (MY ratio) showing how cohort size drives P/E ratios. Abel demonstrates that bequests don't prevent the demographic price cycle. Brooks raises the structural-change confound.

**Key terms:** MY ratio (Middle-aged to Young), projected asset demand, exchange economy vs. production economy, favored cohort effect, price levels vs. returns, general equilibrium.

**Examiner defense strategy:** Don't rely on any single paper. Use GMQ for the mechanism, Poterba for honest engagement with null results, Abel for GE discipline, and Gabaix-Koijen for the empirically identified multiplier. The capstone's

contribution is to model the specific institutional channel that previous papers tested only indirectly.

**What comes next:** Chapter 5 introduces the Gabaix-Koijen framework — the inelastic markets hypothesis that provides the empirical multiplier for translating demographic flows into price impacts.

---

# Chapter 5: The Inelastic Markets Revolution

## Why This Chapter Is the Heart of Everything

Chapters 1-4 built the stage: financial markets where prices are set by supply and demand, a massive demographic cohort cycling through save-then-spend, institutional machinery that channels these flows mechanically, and an academic debate about whether any of it matters for prices.

Chapter 5 answers the "does it matter" question with a resounding yes — and provides the precise quantitative tool the capstone needs. Gabaix and Koijen's 2024 *Econometrica* paper, "In Search of the Origins of Financial Fluctuations: The Inelastic Markets Hypothesis," is not just another contribution to the demographic debate. It is a new framework for understanding stock market fluctuations that happens to be exactly what we need.

Their central finding: **every $1 invested in the stock market increases aggregate market capitalization by approximately $5.** This "multiplier" of ~5× is what converts the dollar flows from our demographic model (Modules 1-4) into predicted market capitalization changes (Module 5).

Let us build the full picture from scratch.

Before diving into their model, it helps to know that this result did not emerge from a vacuum. Gabaix and Koijen systematically review multiplier estimates from across the existing literature (their Table 1 and Figure 2), organized by level of aggregation. At the **micro level** (individual stock price impact), studies of index inclusion events and mutual fund flows find multipliers of about 0.3 to 2.5 — buying one stock pushes its price up, but other stocks can substitute. At the **factor level** (buying a basket like "small caps" or "value stocks"), multiplier estimates jump to around 4.8 to 5.7 from

Morningstar rating changes and fund flow studies. At the **macro level** (aggregate stock market), existing estimates range from 1.5 to 6.5 from dividend payout studies in the U.S. and pension fund rebalancing in Chile. Gabaix and Koijen's own GIV estimate of ~5 falls squarely in this range. The pattern is important: **the multiplier increases with the level of aggregation**, because stocks are close substitutes for each other but stocks and bonds are poor substitutes. Demographic flows are aggregate flows — they enter or exit the total stock market — and therefore face the largest multiplier tier.

## What Does "Inelastic" Mean in This Context?

In Chapter 1 we previewed elasticity. Now let us be precise.

In a stock market context, the **price elasticity of demand** measures how much the aggregate demand for stocks responds to a change in stock prices. Specifically, if the stock price rises by 1%, by what percentage does the quantity of stocks demanded fall?

In a traditional efficient-markets model, the answer is: a lot. There are hordes of rational investors constantly evaluating whether stocks are overvalued or undervalued. If prices rise above fundamental value, these rational investors sell (reducing demand) until the price comes back down. If prices fall below fundamental value, they buy (increasing demand) until the price recovers. This creates high elasticity — prices are tightly anchored to fundamentals by the actions of price-sensitive investors.

Gabaix and Koijen argue that this picture is wrong for the modern U.S. stock market. The reason is simple: **most of the money in the stock market is held by investors who do not respond to prices.**

- **Index funds** (like Vanguard's S&P 500 fund) hold stocks in proportion to the index, regardless of whether any individual stock or the market as a whole is overvalued or undervalued.
- **Target-date funds** maintain a fixed equity allocation based on the glide path, regardless of market valuations.
- **Pension funds and endowments** often follow fixed policy allocations (e.g., 60% equity / 40% bonds) and rebalance to maintain them.
- **401(k) participants** who were auto-enrolled into a default fund and never changed it.

None of these investors provide elasticity. They do not sell when prices are high or buy when prices are low — they maintain fixed allocations regardless. Gabaix and Koijen call this behavior following a "mandate" to hold a fraction θ of wealth in equities.

The investors who *do* provide elasticity — active hedge funds, proprietary trading desks, individual stock pickers — are a small and shrinking share of the market. Gabaix and Koijen's Flow of Funds analysis is damning here: **broker dealers hold less than 0.5% of the equity market** — far too small to absorb large aggregate flows. Hedge funds are slightly larger (~4% before the financial crisis) but, as documented by Ben-David, Franzoni, and Moussawi (2012), they *amplified* selling during 2007-2009 rather than buying the dip — forced by client redemptions and leverage constraints to sell an average of 3.06% of their equity holdings per quarter. When Gabaix and Koijen surveyed economists about which sector provides elasticity to the market, the two most common answers were hedge funds and broker dealers — the very two sectors that are too small or too constrained to play this role. As we will see in Chapter 6, Chinco and Sammon (2024) show that the true passive share is approximately 33.5%, roughly double the commonly cited figure of 16%. When you add quasi-passive investors (closet indexers, balanced funds with fixed allocations), the price-insensitive share of the market is even larger.

The result: aggregate demand for stocks is very price-insensitive — very inelastic. When flows push prices up, there are not enough price-sensitive sellers to push them back down quickly. When flows push prices down, there are not enough bargain hunters to push them back up quickly.

> ***Key Takeaway:*** *"Inelastic" means that prices have to move a lot to restore equilibrium after a flow shock, because there are not enough price-responsive investors to absorb the imbalance. The market does eventually return to equilibrium — but only through large price adjustments rather than small ones.*

## The Undergraduate Example: A Two-Fund World

The most brilliant contribution of the Gabaix-Koijen paper is Section 3.1 — a worked example so simple that an undergraduate can follow it, yet so powerful that it delivers the key insight of the entire paper. Let us walk through it carefully.

**Setup:** Imagine a world with only two investment funds.

- **Fund A:** A pure bond fund. It holds only risk-free bonds and never touches stocks.
- **Fund B:** A mixed fund. It has a mandate to always hold 80% of its wealth in stocks and 20% in bonds. Think of this as a simplified TDF or a balanced fund.

The stock market has one stock, initially priced at $1 per share. There are 4 shares outstanding. So total equity market capitalization is $4.

Fund B has total wealth of $5. Following its 80/20 mandate, it holds $4 in stocks (all 4 shares at $1 each) and $1 in bonds.

Fund A holds only bonds — let us say $5 worth.

**The shock:** Now suppose a household moves $1 from Fund A (the bond fund) to Fund B (the mixed fund). This represents an inflow — like a 401(k) contribution flowing into a balanced/TDF fund.

**What happens mechanically:**

1. Fund A loses $1. Since it only holds bonds, it sells $1 of bonds. No stock market impact.
2. Fund B receives $1. Its total wealth is now $6. Its mandate says: hold 80% in equities. So it needs $4.80 in equities (80% × $6). But it currently holds $4 in stocks. It needs to buy $0.80 more in stocks.
3. But who sells the stocks to Fund B? There is no other stock holder — Fund B already holds all the shares. In a real market, the price would have to rise until the existing shareholders are willing to sell, or until Fund B's mandate is satisfied at a higher price.
4. Here is the key: Fund B does not need to buy *more shares*. It needs to hold 80% of its wealth in equities. If the stock price rises, the value of Fund B's existing stock holdings rises, and at some higher price, Fund B's equity share automatically reaches 80% without any additional purchases.

**Solving for the new price:** After the $1 inflow, Fund B has $6 total. It holds 4 shares at some new price P. Its wealth is 4P + (remaining bonds). For the 80/20 mandate to be satisfied:

$4P / (4P + \text{bonds}) = 0.80$

Working through the algebra (which Gabaix and Koijen do in Section 3.1), the new stock price is approximately $1.05 per share. Fund B's stock holdings are now worth 4 × $1.05 = $4.20.

**The multiplier:** The household invested $1 into Fund B. Of that, $0.80 flowed toward equities (because of the 80/20 mandate). The total market capitalization went from $4.00 to $4.20 — an increase of $0.20 per share × 4 shares = a market cap increase of about $0.20... wait, that does not seem right. Let us redo this more carefully.

Actually, with 4 shares now at $1.05, market cap = $4.20. The increase is $0.20. The equity inflow was $0.80. So the multiplier is $4.20 - $4.00 / $0.80 ≈ ... let me give you the clean version.

The precise Gabaix-Koijen derivation with their calibrated numbers gives a multiplier of approximately 5. The intuition is: when the mixed fund (representing all price-insensitive investors) holds a large fraction of the market, even a small inflow pushes up the price of all existing shares. Since there are many existing shares and only a small inflow, the total market cap change is a large multiple of the flow.

Here is the cleanest way to see it: if the mixed fund holds 80% of all financial wealth in equities, then its elasticity parameter $\zeta = 1 - \theta = 1 - 0.80 = 0.20$. The multiplier is $M = 1/\zeta = 1/0.20 = 5$. That is it. **Every $1 of net flow into equities increases market cap by $5.**

> ***Key Takeaway:*** *The multiplier arises because marginal trades set the price for all existing shares. When most investors are price-insensitive (they don't sell as prices rise), even a small inflow pushes up the price of the entire stock of existing shares. The higher the share of price-insensitive investors (higher θ), the larger the multiplier.*
>
> ***Examiner Alert:*** *An examiner might challenge this example as unrealistic — real markets have millions of participants, not two funds. The defense: this is a pedagogical model, not a literal description. The empirical estimation (GIV, discussed next) validates the multiplier in the real market with all its complexity. The two-fund example teaches the intuition; the econometrics delivers the number.*

## Where Does Elasticity $\zeta = 0.2$ Come From?

The undergraduate example shows *why* the multiplier is large in theory. But what is the actual number in real markets? This requires estimation, and it is hard because of a classic identification problem: prices and flows are simultaneously determined by common factors (macroeconomic news, sentiment, monetary policy).

If you naively regressed stock returns on flows, you would get a biased estimate because both variables respond to the same underlying shocks. Gabaix and Koijen solve this using their **Granular Instrumental Variables (GIV)** methodology, which they developed specifically for this paper.

**The GIV intuition:** Among the hundreds of institutional investors in the 13F filing data, each has idiosyncratic reasons for buying or selling. A pension fund might rebalance because of an unusual contribution spike. A mutual fund family might experience outflows in one sector because of a manager scandal. An endowment might shift its allocation due to a new investment committee chair.

These investor-specific shocks are, by construction, uncorrelated with the aggregate macro factors that simultaneously affect prices. Gabaix and Koijen use them as instruments: they extract the idiosyncratic component of each investor's demand change, take the size-weighted sum (because bigger investors have bigger market impact), and use this as an instrument for aggregate flow.

**Technical details (accessible version):** For each quarter, they observe how much each institutional investor changed their equity holdings. They decompose each investor's change into: (a) a component correlated with what everyone else was doing (the macro component), and (b) a component unique to that investor (the idiosyncratic component). The GIV is the size-weighted sum of the idiosyncratic components.

The trick is that size-weighting minus equal-weighting cancels out the common shocks. A large investor's idiosyncratic shock moves the aggregate market; a small investor's does not. By focusing on the differential between size-weighted and equal-weighted demand changes, they isolate the causal flow-to-price channel.

**The result:** $M \approx 5$, or equivalently $\zeta \approx 0.2$. This means $1 of net inflow causes approximately $5 of market cap increase. Conversely, $1 of net outflow causes approximately $5 of market cap *decrease*.

**How surprising is this?** Very. Before releasing their paper, Gabaix and Koijen surveyed 294 economists and finance professionals with the question: "If a fund buys $1 billion worth of US equities (permanently; it sells bonds to finance that position), slowly over a quarter, how much does the aggregate market value of equities change?"

The median response was M = 0 — surveyed economists believed flows have *zero* price impact, consistent with the textbook efficient-markets model. Among those who gave a positive answer, the median was M = 0.01. Only 2.5% of Twitter respondents and 4% of finance seminar participants answered M ≥ 1. The profession's priors were off by a factor of 500 or more. This underscores how revolutionary the result is — and why examiners trained in the efficient-markets tradition may be instinctively skeptical.

> *Simplified here — the formal GIV econometrics involves precision-weighted cross-sectional demeaning, principal components to extract common factors, and an IV regression. If you want, we can go deeper into the formal econometrics, the first-stage F-statistics, or the robustness checks.*

## Flows Explain Most of Stock Market Fluctuations

One of the paper's most striking results is that flow-driven volatility accounts for the **majority of quarterly stock return variance**. This is shown in their Table 4 and Figure 3.

They construct a measure of "total equity flow" by combining: (a) the aggregate investor flow from their screened 13F data, and (b) corporate share buybacks (which are economically equivalent to a flow into the equity market — the corporation is buying its own shares, reducing supply).

Using this flow measure, they find:

**Flows and returns co-move strongly.** The left panel of Figure 3 shows quarterly stock market returns plotted against quarterly net flows. The relationship is steep and positive — quarters with high inflows have high returns, and quarters with low or negative flows have poor returns.

**Survey beliefs predict flows predict returns.** Using Gallup survey data on investor expectations, they show that optimistic beliefs predict inflows (people who expect stocks to rise put money in), and these inflows predict returns (because the multiplier amplifies them). Survey expectations alone explain about 23% of quarterly return variance. Combined with GDP and other fundamentals, their model explains about 58%.

**The causal chain is:** Beliefs/sentiment → Flows → Price changes (via multiplier). This is profoundly different from the traditional story (Fundamentals → Price changes). It does not mean fundamentals are irrelevant — dividend changes (δ) do affect prices. But the flow channel is empirically larger.

***Connection to Our Model:*** *If flows explain most of short-run market fluctuations, then demographic flows — which we can project forward for 20 years — should explain a significant portion of long-run price-level changes. This is the bridge from Module 4 (net flows) to Module 5 (price impact). We are not predicting quarterly returns — we are projecting the cumulative effect of sustained net outflows on market capitalization over decades.*

## The General Equilibrium Model: Equity Premium, Predictability, and the Pricing Kernel

Gabaix and Koijen do not stop at estimating the multiplier. In Section 5, they build a full general equilibrium model to show that inelastic markets can explain several deep puzzles in finance. For the capstone, three results matter:

**Result 1: The equity premium puzzle dissolves.** The equity premium — the ~6% annual excess return stocks earn over bonds — has been notoriously difficult to explain with standard models. The traditional story requires implausibly high risk aversion ($\gamma \approx 50$ or more). Gabaix and Koijen show that with inelastic markets and $\gamma = 2$ (a perfectly reasonable value), the model generates the correct equity premium. The key insight: because flow shocks are amplified by the multiplier, stock return volatility is much higher than dividend volatility. Investors demand compensation for this flow-driven volatility, generating a large premium.

The formula is: $\bar{\pi} = \gamma \times \sigma^2_r$, where $\bar{\pi}$ is the equity premium, $\gamma$ is risk aversion, and $\sigma^2_r$ is the variance of stock returns. With $\gamma = 2$ and flow shocks generating ~90% of return variance, this matches the data without exotic assumptions.

**Result 2: Return predictability.** In the model, equity premium varies over time: $\pi_t = \bar{\pi} + b \times f_t$. After a period of high inflows, prices are elevated and the equity premium is low (stocks are expensive relative to fundamentals). After a period of outflows, prices are depressed and the equity premium is high (stocks are cheap). This generates the well-documented predictability of stock returns by the price-dividend ratio — high P/D predicts low future returns, not because markets are irrational, but because they are inelastic.

**Result 3: The flow-based pricing kernel.** Traditionally, asset pricing models use a "stochastic discount factor" (SDF) based on marginal utility of consumption. But consumption-based SDFs have well-known empirical failures. Gabaix and Koijen show that in their model, the SDF should be based on flows, not consumption. Consumption

explains bond pricing (which is roughly risk-free) but fails for equities. Flows explain equity pricing. The SDF is a "symptom, not a cause" — it records the influence of flows on prices.

> ***Key Takeaway for the capstone:*** *The GE model tells us that the multiplier is not just a short-run phenomenon. It has equilibrium implications: sustained outflows should permanently reduce price levels and permanently increase the equity premium. This means demographic outflows do not just create a temporary dip — they create a prolonged period of higher expected returns (good for future buyers, bad for current holders). This is exactly the "favored cohort effect" from Geanakoplos, now with a formal mechanism.*
>
> ***Examiner Alert:*** *An examiner might argue that the GE model's assumption of boundedly rational consumers (who do not undo fund mandates) is strong. The defense: Giglio et al. (2021) provide direct evidence. Using Vanguard data on millions of investors, they show that even when beliefs change dramatically (e.g., after COVID), portfolio allocations barely move. Beliefs are volatile; portfolios are sticky. This is exactly the behavioral assumption Gabaix-Koijen need.*

## Corporate Buybacks: The Key Counterforce

Section 6.2 of Gabaix-Koijen introduces a critically important element for the capstone: corporate share buybacks as an offsetting flow.

**What buybacks are:** When a corporation uses its cash (from profits or borrowing) to buy back its own shares on the open market, it is economically equivalent to an inflow into the equity market. The company is a buyer of equity, reducing the supply of outstanding shares. Through the multiplier, buybacks increase stock prices.

**Why this matters for the capstone:** If baby boomer retirements generate net outflows of $X billion per year, but corporate buybacks generate inflows of $Y billion per year, the net market-moving flow is only $(X - Y) billion. Buybacks partially or fully offset demographic selling pressure.

**The numbers are enormous.** U.S. corporations have been buying back hundreds of billions of dollars of stock annually. S&P 500 buybacks have exceeded $800 billion in recent years. This is a significant counterforce.

**But there is a catch.** Buybacks are procyclical — companies buy back more stock when profits are high and the economy is strong, and cut buybacks when times are tough. If demographic outflows depress stock prices and weaken the economy,

corporate profits may fall, reducing buybacks precisely when they are most needed as a counterforce. This creates a potential positive feedback loop: outflows depress prices → lower profits → fewer buybacks → less offsetting inflow → further price depression.

Gabaix and Koijen note that in their model, if firms in the aggregate buy back $1 worth of equity, that can increase aggregate valuations by approximately $5 (the same multiplier applies). Buybacks are not neutral under Modigliani-Miller because the market is inelastic — the traditional irrelevance theorem breaks down when demand is price-insensitive.

> ***Connection to Our Model:*** *Buybacks are the single most important counterforce to demographic outflows. The capstone should include buyback projections as a key scenario variable. Conservative scenario: buybacks decline as outflows mount. Moderate: buybacks remain at historical levels. Aggressive: buybacks increase (perhaps driven by companies taking advantage of lower valuations). The net flow — demographic outflows minus buybacks — is what the multiplier should be applied to.*
>
> ***Examiner Alert:*** *An examiner could argue that buybacks will fully offset demographic outflows, rendering the thesis irrelevant. The response: buybacks are discretionary and procyclical — they cannot be relied upon to mechanically offset a structural, mandatory outflow like RMDs. Moreover, buybacks funded by debt add leverage, which increases financial fragility. The capstone should present buybacks as a moderating force, not an elimination of risk.*

## Government "Equity QE" — A Policy Digression

Gabaix-Koijen's Section 6.1 briefly considers what happens when a government buys equities directly — "quantitative easing for stocks." If the government purchases 1% of the stock market, the multiplier predicts a ~5% price increase.

They cite real-world examples: Hong Kong's government bought ~6% of the market during the 1998 Asian financial crisis, and the market rallied ~24% — approximately a 4x multiplier, consistent with their estimates. The Bank of Japan has been buying ETFs for years. China's "national team" intervenes during market stress.

For the capstone, this is a policy implication rather than a model input. If demographic outflows become severe enough to threaten market stability, government intervention is a possibility — one that operates through exactly the mechanism we are modeling.

## What About Slow-Moving Capital?

One question the inelastic markets hypothesis raises: if prices deviate significantly from fundamentals (pushed up or down by flows), why don't rational arbitrageurs rush in to profit?

Mitchell, Pedersen, and Pulvino (2007), in their paper "Slow Moving Capital," provide the answer. Capital moves slowly for several structural reasons:

**Institutional constraints:** Most investment capital is locked in mandates. A pension fund with a 60/40 policy allocation cannot suddenly shift to 80% equities even if stocks become very cheap. The mandate is the mandate.

**Career risk:** A fund manager who buys aggressively into a falling market risks underperforming benchmarks in the short run. If performance lags for two or three quarters, investors may pull their money (outflows from the fund), forcing the manager to sell at the worst possible time. This creates a perverse incentive: even managers who believe stocks are cheap may avoid buying because they cannot survive the short-term drawdown.

**Leverage constraints:** Sophisticated arbitrageurs who recognize mispricings may face margin calls if they are leveraged. A price decline beyond what they can fund forces them to unwind their positions, exacerbating the very mispricing they were trying to correct.

**Information uncertainty:** During periods of market stress, even sophisticated investors are uncertain whether price declines reflect fundamentals or flow-driven overshooting. They wait for clarity, which delays their entry.

The result is that flow-driven price displacements persist for months or years, not hours or days. Mitchell et al. document this across several market episodes, showing that capital takes weeks to months to flow to opportunities.

> ***Connection to Our Model:*** *Slow-moving capital tells us that the price impact of demographic outflows will not be immediately offset by rational buyers. Even if prices fall below fundamental value, the correction will take time. This means the capstone's projected price impacts should be interpreted as persistent, not transient.*

## ETFs and Volatility Amplification

Ben-David, Franzoni, and Moussawi (2018), "Do ETFs Increase Volatility?" in the *Journal of Finance*, provide complementary evidence that the rise of passive investing amplifies price movements.

Their core finding: stocks with higher ETF ownership experience greater non-fundamental volatility. The mechanism is that ETF creation and redemption (the process by which ETF shares are created or destroyed to match demand) transmits flows directly to underlying stocks. When ETF investors panic and sell, the ETF market maker redeems shares and sells the underlying stocks — mechanically. When ETF investors are euphoric and buy, the reverse happens.

Because ETF flows operate at the basket level (buying or selling the entire index), they create correlated price movements across all constituent stocks, regardless of individual company fundamentals. This increases co-movement and aggregate volatility.

For the capstone, this finding reinforces the inelasticity story: the same passive vehicles that make the market inelastic (index funds, ETFs) also amplify the price impact of flows. Demographic outflows flowing through ETF redemptions will create correlated selling pressure across all index constituents.

## The True Passive Share: Chinco and Sammon (2024)

How much of the stock market is actually held by passive investors? The commonly cited figure from ICI (Investment Company Institute) data is roughly 16% in index mutual funds and ETFs. But Chinco and Sammon (2024) show this dramatically understates the truth.

Their paper, "The Passive Ownership Share Is Double What You Think It Is," published in the *Journal of Financial Economics*, demonstrates that the true passive share is approximately **33.5%** of U.S. equities. Their methodology is clever: each time a stock is added to or dropped from a major index, they ask "How much money would have to be tracking that index to explain the huge spike in rebalancing volume we observe on reconstitution day?" The answer consistently exceeds what reported index fund holdings would predict. The discrepancy arises because:

**Index mutual funds and ETFs are only part of the picture.** Many "passive" holdings exist inside institutional vehicles that are not tracked by standard ICI data — including separate accounts that track indices, collective investment trusts (CITs), and internal index strategies used by pension funds and endowments.

**Closet indexing adds more.** Many funds marketed as "active" actually hold portfolios that closely track benchmark indices. Their active share is low, meaning they are effectively passive for price-discovery purposes.

Chinco and Sammon also show that passive investors rebalance mechanically when index composition changes, creating predictable price pressure around reconstitution events (when stocks are added to or dropped from major indices).

For the capstone, the Chinco-Sammon finding means the market is more inelastic than even Gabaix-Koijen's baseline calibration suggests. If the passive share is 33.5% rather than 16%, and growing at roughly 1-1.5 percentage points per year, the multiplier is likely increasing over time. This is exactly the time-varying multiplier hypothesis we use in Module 5's moderate and aggressive scenarios.

## Beliefs and Portfolios: Giglio et al. (2021)

One potential escape from the inelasticity story: what if investors rationally adjust their portfolios in response to changing conditions? If prices fall due to demographic outflows, maybe investors become more optimistic (recognizing the value opportunity) and increase their equity allocations, offsetting the outflows.

Giglio, Maggiori, Stroebel, and Utkus (2021), "Five Facts About Beliefs and Portfolios" in the *American Economic Review*, test this directly using proprietary data from Vanguard on millions of individual investors' survey responses (about market expectations) and actual portfolio allocations.

Their five facts include the devastating finding that **beliefs are volatile but portfolios barely move.** Specifically:

- Investors change their expectations about market returns dramatically over time (in response to COVID, market crashes, elections, etc.)
- But these belief changes translate into negligible portfolio adjustments
- The correlation between belief changes and portfolio changes is extremely low
- Even investors who dramatically revise their expected returns downward do not meaningfully reduce their equity allocations

This is the micro-level evidence that supports the macro-level inelasticity finding. People do not rebalance in response to changing conditions — not because they cannot,

but because of inertia, complexity, and the effort required. This is the same "power of suggestion" mechanism Madrian and Shea identified, now observed at massive scale.

For the capstone, Giglio et al. answer a critical objection: "won't investors respond to falling prices by buying more?" The empirical answer is: no, they will not — at least not nearly enough to offset the mechanical outflows from RMDs and TDF de-risking.

## Chapter 5 Summary

**Core concepts:** Markets are inelastic because price-insensitive investors (passive funds, TDFs, pension funds) dominate holdings. The multiplier is ~5× — $1 of flow changes market cap by ~$5. Flows explain most of short-run return variance. The GE model explains equity premium, return predictability, and pricing kernel through flows.

**Key terms and parameters:** Elasticity $\zeta \approx 0.2$, Multiplier $M = 1/\zeta \approx 5$, mandate $\theta$ (fraction in equities), GIV (Granular Instrumental Variables), equity premium $\bar{\pi} = \gamma \times \sigma^2_r$, flow-based SDF.

**Papers drawn from:** Gabaix & Koijen (2024) — the entire framework; Chinco & Sammon (2024) — true passive share; Giglio et al. (2021) — beliefs vs. portfolios; Mitchell et al. (2007) — slow-moving capital; Ben-David et al. (2018) — ETF volatility amplification.

**Why it matters for the capstone:** The multiplier is the quantitative bridge from Module 4 (net flows) to Module 5 (price impact). Without the multiplier, we can only project flows. With it, we can project market capitalization changes.

**What comes next:** Chapter 6 asks whether the multiplier is changing over time — and introduces a second, independent method for measuring market inelasticity.

---

# Chapter 6: How Competitive Is the Stock Market Really?

## Why We Need a Second Paper

Gabaix and Koijen estimate the multiplier at the **aggregate** level — one number for the entire stock market, based on time-series variation. Their approach answers: "when more money flows into the stock market as a whole, how much does the total market cap change?"

But for the capstone, we also need to know: **is the market becoming more or less inelastic over time?** If passive investing continues to grow (as all trends suggest), does the multiplier increase? And is there independent confirmation that the Gabaix-Koijen result is not an artifact of their particular methodology?

Haddad, Huebner, and Loualiche (2025), "How Competitive Is the Stock Market?" in the *American Economic Review*, provides both. They approach the inelasticity question from a completely different angle — the **cross-section** of individual stocks rather than the time-series of the aggregate market — and reach the same qualitative conclusion through an independent methodology. They also provide a formal framework for answering how inelasticity changes as passive investing grows.

## The Demand Systems Approach

Haddad et al. estimate a **demand system** — a model of how each institutional investor's demand for each individual stock responds to changes in that stock's price. This is a fundamentally different level of granularity from Gabaix-Koijen.

Think of it this way. Gabaix-Koijen ask: "if $1 flows into the stock market, how much does total market cap change?" Haddad et al. ask: "if the price of *Apple specifically* rises by 1%, how do *each of the thousands of institutional investors* adjust their Apple holdings?"

By estimating these stock-level demand functions for the universe of institutional investors, they can aggregate up to compute the market-level elasticity — but they can also decompose it by investor type, by stock, and over time.

Their specification builds on the demand system framework of Koijen and Yogo (2019), but with a crucial innovation: **strategic responses.** In their model, when one investor becomes less elastic (say, by switching to passive indexing), other investors can respond by becoming *more* elastic (trading more aggressively to exploit the mispricings created by the passive investor's rigidity).

## The Strategic Substitutability Result: $\chi = 3$

This is the paper's most important finding for the capstone. Haddad et al. estimate the **strategic substitutability parameter $\chi \approx 3$**, which measures how much active investors increase their elasticity when passive investing grows.

Here is what $\chi = 3$ means in plain language. When the passive share of the market increases, some elasticity is "lost" (passive investors provide zero elasticity). But active investors observe that prices are now more dislocated from fundamentals (because there are fewer elastic investors keeping prices anchored). Seeing bigger mispricings, active investors trade more aggressively — they provide *more* elasticity per investor.

The question is: how much of the lost elasticity do active investors replace? The answer depends on $\chi$:

- If $\chi = 0$: No strategic response. Active investors do not change their behavior. All lost elasticity from passive growth is truly lost. The **pass-through** from passive growth to inelasticity is 100%.
- If $\chi \to \infty$: Perfect strategic response. Active investors fully compensate. No matter how much passive investing grows, aggregate elasticity stays the same. Pass-through is 0%.
- **$\chi = 3$ (the empirical estimate):** Partial strategic response. Active investors step up, but not enough to fully offset the passive share increase. The **pass-through is approximately 1/3** — meaning one-third of the elasticity lost to passive investing is truly lost (not compensated by active traders).

## The Pass-Through Formula

Haddad et al. derive a simple, powerful formula for the pass-through from a change in passive share to a change in aggregate elasticity:

**Pass-through = $1 / (1 + \chi \times |\text{Active}|)$**

Where $|\text{Active}|$ is the fraction of the market held by active investors and $\chi$ is the strategic substitutability parameter.

With $\chi = 3$ and $|\text{Active}| \approx 0.67$ (passive share ~33%):

Pass-through $= 1 / (1 + 3 \times 0.67) = 1 / (1 + 2.01) = 1 / 3.01 \approx 0.33$

So about 33% of each percentage point increase in passive share translates into reduced elasticity.

**The historical calculation:** Over the past 20 years, the passive share has increased by roughly 32 percentage points (from near zero to about 33.5%, per Chinco and Sammon). With a pass-through of ~0.33:

Elasticity reduction = 0.33 × 32% ≈ 11%

This is the headline result: **the rise in passive investing over the last 20 years has made the demand for individual stocks 11% more inelastic.**

> ***Key Takeaway:*** *Active investors do respond to the growth of passive investing — they trade more aggressively. But their response only offsets about two-thirds of the lost elasticity. The remaining one-third is a permanent reduction in market competitiveness.*
>
> ***Examiner Alert:*** *An examiner might argue that $\chi = 3$ means active investors are doing a pretty good job compensating — two-thirds offset is substantial. The defense: even the remaining one-third matters enormously in dollar terms when applied to a $50+ trillion market. An 11% decrease in elasticity over 20 years means the multiplier has gone from (hypothetically) 5x to about 5.6x. Applied to the hundreds of billions of dollars of demographic flows we project, the difference is trillions of dollars in market cap impact. Moreover, the trend is accelerating — passive share is still growing at 1-1.5 pp/year.*

## Russell Reconstitution: A Natural Experiment

How do Haddad et al. know their estimates are right? They validate using **Russell index reconstitutions** — one of the cleanest natural experiments in finance.

Every June, the Russell 1000 (large caps) and Russell 2000 (small caps) indices are reconstituted based on market capitalization rankings. Stocks near the boundary between the two indices may be shifted from one to the other. When a stock moves from the Russell 2000 to the Russell 1000 (or vice versa), the passive funds tracking each index must buy or sell accordingly. This creates an exogenous flow shock at the stock level — the buying and selling is driven entirely by the index rule, not by any fundamental change in the company.

Haddad et al. use these reconstitution events to test their model's predictions about pass-through. Their model predicts specific price impacts and strategic responses for stocks near the Russell boundary, and these predictions match the data well.

This validation is important for the capstone because it means the pass-through formula is not just a theoretical construct — it generates correct predictions in a controlled setting.

## Two Methods, Same Conclusion

Here is the meta-point that ties Chapters 5 and 6 together, and that an examiner will find compelling:

**Gabaix and Koijen (2024)** approach inelasticity from the **aggregate, time-series** direction. They ask how total flows affect total market returns. They use GIV as an identification strategy. They find $M \approx 5$ ($\zeta \approx 0.2$).

**Haddad, Huebner, and Loualiche (2025)** approach inelasticity from the **cross-sectional, stock-level** direction. They ask how each investor's demand for each stock responds to price changes. They use a demand systems framework with strategic responses. They find that markets are inelastic and becoming more so.

These are independent methodologies, different datasets, different estimation strategies, different journals (*Econometrica* vs. *AER*) — and they reach the same conclusion. When two independent research teams, using fundamentally different approaches, converge on the same finding, the evidence is much stronger than either paper alone.

For the capstone, this convergence provides the intellectual foundation for applying the multiplier with confidence. We are not relying on one paper's estimate; we are relying on a consistent finding across the two leading papers in the field.

## The Time-Varying Multiplier: How We Use Haddad et al. in Module 5

Gabaix-Koijen give us the *level* of the multiplier ($M \approx 5$). Haddad et al. give us the *trend* — how the multiplier changes as passive investing grows. Combining the two creates the time-varying multiplier that our capstone's Module 5 uses.

**The logic:**

1. Start with baseline multiplier $M = 5$ (Gabaix-Koijen, based on historical average)

2. Project future passive share growth (from ICI trend data: ~1-1.5 pp/year increase)
3. For each future year, compute how much passive share has grown since baseline
4. Apply the Haddad et al. pass-through formula to calculate the cumulative elasticity decline
5. Compute the new multiplier: $M_new = 1 / \zeta_new$, where $\zeta_new = \zeta_baseline \times (1 - \text{cumulative elasticity decline})$

**Three scenarios:**

- **Conservative:** Static $M = 5$. Ignore the time trend entirely. This is the lower bound.
- **Moderate:** Time-varying, reaching $M \approx 6$ by 2035 as passive share grows to ~45%.
- **Aggressive:** Time-varying, reaching $M \approx 8$ by 2040 if passive growth accelerates.

  ***Examiner Alert (important caveat):*** *Haddad et al.'s estimates are cross-sectional (individual stock elasticities) while Gabaix-Koijen's are aggregate (whole-market elasticity). These are not the same thing, and directly combining their magnitudes requires assumptions. The capstone should frame this as: "Both papers find low and declining elasticity through different lenses. We use Gabaix-Koijen for the level and Haddad et al. for the trend direction, bracketing magnitude uncertainty through scenarios." This is an honest limitation, not a flaw — it is the appropriate way to handle methodological incompatibility between two strong papers.*

## Chapter 6 Summary

**Core concepts:** Haddad et al. estimate stock-level demand elasticities using a demand systems approach with strategic responses. Active investors partially compensate for passive growth ($\chi = 3$, absorbing two-thirds), but one-third of lost elasticity passes through to reduced market competitiveness. Over 20 years, this produced an 11% reduction in demand elasticity.

**Key terms and parameters:** Demand system, strategic substitutability ($\chi \approx 3$), pass-through ($\approx 1/3$), Russell reconstitution, cross-sectional vs. aggregate elasticity.

**Papers drawn from:** Haddad, Huebner & Loualiche (2025) — the entire framework; Chinco & Sammon (2024) — true passive share for calibration.

**Why it matters for the capstone:** Haddad et al. provide the time-varying dimension — the multiplier is not fixed at 5, it is growing as passive investing expands. This makes the capstone's price impact projections time-dependent and forward-looking. Combined with Gabaix-Koijen's level estimate, we have a complete quantitative toolkit.

**What comes next:** Chapter 7 brings everything together — connecting all six core papers and six complementary papers to the five modules of the capstone model.

---

# Chapter 7: Synthesis — How All the Pieces Connect

## The Intellectual Arc

Let us step back and see how the six chapters we have covered tell one coherent story, and how that story maps to the capstone model.

**The story in one paragraph:** People save when working and spend when retired (Chapter 2). The baby boom created a massive cohort that flooded equity markets with savings during the 1980s-2000s and is now beginning to withdraw (Chapter 2). American retirement institutions — 401(k)s with auto-enrollment, target-date funds, and RMDs — channel these demographic flows mechanically and without price sensitivity (Chapter 3). The academic debate over whether demographics affect markets found weak effects in the pre-passive era (Chapter 4), but the inelastic markets hypothesis reveals that in today's passive-dominated market, flows have amplified price effects — a 5× multiplier (Chapter 5). Independent evidence confirms that markets are inelastic and becoming more so as passive investing grows (Chapter 6). Combining these pieces, the capstone projects the net flow trajectory and applies the multiplier to estimate market capitalization impacts.

**The story in one sentence:** The same institutional machinery that amplified inflows during the saving phase will amplify outflows during the retirement phase.

## Mapping Papers to Modules

Here is how each of the twelve papers (six core, six complementary) connects to the five modules of the capstone model:

### Module 1: Demographics (U.S. Population by Age Cohort, 2025-2045)

**Primary source:** Census Bureau National Population Projections (data, not a paper).

**Papers that provide context:** - Geanakoplos, Magill & Quinzii (2004) — the MY ratio framework shows *why* age structure matters for equity demand. Provides the intellectual justification for building a cohort-based model. - Abel (2001) — confirms that the demographic price cycle is robust to bequest behavior, meaning cohort-based flow projections do not need to be adjusted for inheritance effects.

### Module 2: Inflows (Annual Equity Contributions by Age Cohort)

**Key parameters:** - Participation rates by age → Madrian & Shea (2001): auto-enrollment drives ~85% participation - Contribution rates by age → Vanguard "How America Saves" data, Madrian & Shea default stickiness - TDF equity allocation by age (glide path) → Parker, Schoar & Sun (2020): 90% equity at 25, declining to 30-40% past retirement - Auto-enrollment mandate → SECURE 2.0 provisions

**Papers drawn from:** - Madrian & Shea (2001) — the behavioral foundation: defaults dominate, participation is high, allocations are sticky - Parker, Schoar & Sun (2020) — TDF growth, glide path mechanics, rebalancing speed, macrocontrarian behavior - Giglio et al. (2021) — confirms portfolio inertia at massive scale: beliefs change, allocations do not

### Module 3: Outflows (Annual Equity Withdrawals by Age Cohort)

**Key parameters:** - RMD schedule → IRS Uniform Lifetime Table, SECURE 2.0 age thresholds (73 or 75) - Voluntary withdrawal rates by age → Poterba, Venti & Wise (drawdown data), HRS microdata - TDF de-risking → Parker et al. (2020): glide path reduces equity allocation ~1-2 percentage points per year in the decade before and after retirement

**Papers drawn from:** - Poterba, Venti & Wise — actual withdrawal patterns and wealth composition at retirement - Parker, Schoar & Sun (2020) — TDF de-risking as a mechanical outflow channel - SECURE 2.0 legislation — RMD age changes

### Module 4: Net Flow Aggregation

This is the pure accounting module: sum inflows (Module 2) minus outflows (Module 3) across all cohorts for each year, using Monte Carlo simulation to propagate uncertainty through demographic projections, participation rate uncertainty, return uncertainty, and behavioral parameter uncertainty.

**Papers that inform uncertainty ranges:** - Poterba (2001) — his null result for returns reminds us that demographic effects on returns are uncertain, which we should reflect in wide confidence bands - Geanakoplos et al. (2004) — the MY ratio trajectory provides an independent sanity check on our net flow timing

### Module 5: Price Impact Application

**Key formula:** $\Delta$Market Cap = M × Net Flow, where $M = 1/\zeta$

**Baseline multiplier:** $M \approx 5$ ($\zeta \approx 0.2$) from Gabaix & Koijen (2024)

**Time-varying multiplier:** Uses Haddad et al. (2025) pass-through formula to project how M changes as passive share grows: - Conservative: M = 5 (static) - Moderate: M grows from 5 to ~6 by 2035 - Aggressive: M grows from 5 to ~8 by 2040

**Papers drawn from:** - Gabaix & Koijen (2024) — the multiplier level, GIV estimation, GE model implications - Haddad, Huebner & Loualiche (2025) — the time trend in elasticity, pass-through formula, $\chi = 3$ - Chinco & Sammon (2024) — true passive share for calibrating the time-varying multiplier - Ben-David, Franzoni & Moussawi (2018) — ETF amplification mechanism - Mitchell, Pedersen & Pulvino (2007) — slow-moving capital explains why price impacts persist

## The Cross-Paper Connections That Make This Work

The capstone is not just a collection of papers stapled together. It works because the papers connect to each other in ways that make the argument stronger than any single paper alone.

**Connection 1: Defaults → Inelasticity.** Madrian & Shea show that auto-enrollment creates massive price-insensitive flows. Gabaix-Koijen show that price-insensitive flows make markets inelastic. Therefore: auto-enrollment is a root cause of market inelasticity. The same policy that secured millions of Americans' retirement savings also made equity markets less elastic and more susceptible to flow-driven price movements.

**Connection 2: Two methods, same conclusion.** Gabaix-Koijen (aggregate, time-series, GIV) and Haddad et al. (cross-sectional, demand system, Russell reconstitution) independently confirm low and declining elasticity. Neither depends on the other's methodology. This convergence is powerful evidence.

**Connection 3: Demographics drive levels, not returns.** Poterba finds demographics do not predict returns but do relate to price levels. Gabaix-Koijen's multiplier operates on price levels, not returns. These are consistent: demographic flows change the level of market capitalization (through the multiplier), but the market may price this in efficiently, producing no abnormal return. The capstone's contribution is projecting the *level* change, not predicting *returns*.

**Connection 4: Active investors partially compensate.** Haddad et al. show active investors absorb two-thirds of the elasticity gap created by passive growth. This means demographic outflows will be partially offset by active buying (value investors stepping in as prices fall). But one-third passes through — the outflow impact is real, just moderated.

**Connection 5: The correction is slow.** Mitchell et al. show capital moves slowly. Even if demographic outflows create mispricings (prices below fundamental value), the correction will take months to years. This means the capstone's projected price impacts should be interpreted as persistent effects, not short-lived blips.

**Connection 6: Bequests do not save the market.** Abel shows that even with full bequests, the demographic price cycle persists. This removes a potential objection. Combined with Gabaix-Koijen's multiplier, it means: demographic outflows will produce persistent, amplified price effects regardless of inheritance behavior.

**Connection 7: International flows will not offset.** Geanakoplos et al. note that foreign investors could potentially offset domestic demographic pressure. But the largest foreign holders of U.S. equities (Europe, Japan, Korea, China) face the same or worse demographic transitions. If anything, foreign outflows will reinforce rather than offset U.S. boomer selling pressure. This turns a potential limitation into a strength: the capstone provides a *lower bound* estimate because it only models U.S. domestic flows.

## The Cascading Problems of Inelasticity

This section explores the broader implications of the inelasticity story — what happens beyond the immediate portfolio impact when markets become less elastic and

demographic outflows create sustained price depression. This is the "so what" beyond the numbers.

**Step 1: Prices detach from fundamentals.** In an inelastic market, prices are driven more by flows and less by fundamental value. If demographic outflows push prices below fundamental value, the price signal is "wrong" — it does not accurately reflect the present value of future cash flows.

**Step 2: Wrong price signals → mispriced cost of capital.** Companies use stock prices to determine their cost of equity capital. If stock prices are artificially depressed by demographic outflows (not by deteriorating business fundamentals), companies face a higher cost of capital than warranted. This makes it more expensive for firms to raise equity financing, invest in projects, and expand.

**Step 3: Mispriced cost of capital → distorted investment.** Higher-than-warranted cost of capital means some economically productive projects are not funded. Companies underinvest relative to what the economy needs. This reduces future productivity growth.

**Step 4: Reduced investment → lower GDP growth.** If the corporate sector underinvests because of flow-driven price depression, the real economy grows more slowly. This is how a financial market distortion becomes a real economic problem.

**Step 5: Potential for rapid correction (both directions).** When prices are held away from fundamentals by sustained flows, the eventual correction can be rapid and destabilizing. If something triggers a sudden reversion (for example, a policy change, a foreign capital influx, or a shift in domestic saving behavior), prices could snap back violently. In an inelastic market, *upward* flow shocks also produce outsized price movements — the multiplier works in both directions.

> ***Examiner Alert:*** *The cascading problems argument extends well beyond what any single paper proves empirically. Steps 1-2 are well-supported (Gabaix-Koijen formalize the flow-price channel; standard corporate finance establishes cost-of-capital effects). Steps 3-4 are theoretical plausible but empirically untested in the specific demographic context. Step 5 is speculative. The capstone should present this as a "plausible implications" discussion, not as a proven chain. Label each step with its evidentiary status.*

## Corporate Buybacks as Counterforce: A Detailed Treatment

We introduced buybacks in Chapter 5, but they deserve deeper treatment here because they are the single most important force that could moderate the capstone's projected outflow impacts.

**The mechanism:** When a corporation buys back its own shares, it is economically equivalent to an inflow into the equity market. Buybacks reduce the supply of outstanding shares and, through the Gabaix-Koijen multiplier, increase stock prices. In recent years, S&P 500 buybacks have exceeded $800 billion annually.

**Why buybacks matter for the capstone:** If annual demographic outflows reach (hypothetically) $200 billion and annual buybacks are $800 billion, the net flow is still +$600 billion — positive, meaning the market is still receiving net inflows. Demographic outflows would reduce the *growth rate* of net inflows rather than turning net flows negative. The crossover point — when demographic outflows exceed the net of all other inflows including buybacks — is later and potentially never reached.

**Why buybacks are not a reliable counterforce:**

First, buybacks are **procyclical**. Companies buy back more stock when profits are high and cash flow is strong. A sustained period of demographic-driven price weakness could reduce corporate profits (through wealth effects and reduced consumer spending), which would reduce buybacks precisely when they are most needed.

Second, buybacks are **discretionary**. Management can stop buybacks at any time. During COVID-19, many companies suspended buyback programs. During the 2008-2009 financial crisis, buybacks collapsed. RMDs, by contrast, are mandatory. A counterforce that disappears during stress is unreliable.

Third, buybacks funded by **debt** increase corporate leverage, which makes the financial system more fragile. If buybacks are funded by profits, they are sustainable. If funded by cheap debt, they are vulnerable to rising interest rates.

Fourth, there is **growing political and regulatory pressure** against buybacks. Legislation has been proposed to restrict or tax buybacks more heavily. The 2022 Inflation Reduction Act introduced a 1% excise tax on buybacks. Future legislation could reduce buyback volumes.

**How to handle in the model:** Include buyback projections as a key scenario variable.

- **Conservative (worst case for thesis):** Buybacks remain at current high levels, partially offsetting demographic outflows. Net flows turn negative later.

- **Moderate:** Buybacks remain roughly constant in nominal terms but decline as a share of market cap. Demographic outflows eventually dominate.
- **Aggressive (worst case for markets):** Buybacks decline as profits fall in response to demographic-driven price weakness. Feedback loop amplifies outflow impact.

## The Value Investing Paradox

The inelastic markets framework creates a paradox for value investors that is worth understanding.

If demographic outflows push prices below fundamental value (through the mechanical selling of index funds), this creates a textbook value investing opportunity. Stocks are cheap relative to their intrinsic worth. A rational, patient investor should buy.

But here is the problem: in an inelastic market, there are *not enough* of these rational buyers to fully offset the selling pressure. That is what inelasticity *means* — insufficient price-sensitive demand. Haddad et al. confirm this: active investors do step up ($\chi = 3$), but they only absorb two-thirds of the gap.

Moreover, value investing in an inelastic market is especially dangerous because: - **The mispricing can persist for years** (Mitchell et al.'s slow-moving capital) - **The drawdown before recovery can be severe** (active investors face career risk and leverage constraints) - **Most active funds have investment horizons of 1-3 years**, too short to wait out a decade-long demographic wave

The implication: value investing becomes theoretically *more* profitable in inelastic markets (bigger mispricings) but practically *harder* to execute (longer holding periods, deeper drawdowns, career risk). This means demographic outflows will not be fully arbitraged away by sophisticated investors. The price impact will persist.

## International Synchronization: Why Foreign Investors Will Not Save Us

A natural objection: "The U.S. is not a closed economy. Foreign investors can step in when American retirees sell."

This objection fails for two reasons:

**Reason 1: Synchronized aging.** The countries that hold the most U.S. equities are themselves facing severe demographic transitions: - **Japan** — decades ahead of the U.S. in aging; working-age population shrinking since the 1990s - **Europe** — lower fertility than the U.S.; Germany's old-age dependency ratio already exceeds the U.S.'s - **South Korea** — the lowest fertility rate in the world (below 0.8 children per woman) - **China** — working-age population began shrinking in the 2010s; one-child policy consequences hitting now

These countries' pension systems and sovereign wealth funds are the largest foreign holders of U.S. equities. As their own populations age, they will face the same drawdown pressures — not just domestically, but in their U.S. holdings.

**Reason 2: Currency risk.** Foreign investors buying U.S. equities bear currency risk. If U.S. equity prices are declining due to domestic outflows, the dollar may weaken simultaneously (less demand for dollar-denominated assets). This creates a double headwind for foreign buyers: declining equity prices *and* currency losses. This reduces the incentive for foreign investors to step in.

Geanakoplos et al. explicitly flag this issue: "the future path of U.S. equity prices may well depend on the joint demography of countries participating on the U.S. equity market." Our capstone models only U.S. domestic flows, which means our projections are a **lower bound** on total demographic selling pressure. If foreign outflows are correlated (as all evidence suggests), the actual impact is larger than we project. This turns a limitation into a strength.

## Big-Cap Concentration Effect

Here is a subtlety that examiners with finance knowledge will appreciate:

Demographic outflows flow through index funds, which sell stocks in proportion to index weights. The largest companies (Apple, Microsoft, Nvidia, Amazon) have the highest index weights and therefore receive the most selling pressure in dollar terms.

Per Haddad et al.'s framework, these same large-cap stocks have the **highest passive ownership** (because index funds hold them by construction) and therefore the **lowest demand elasticity**. This means the same dollar of selling creates a larger price impact on big-cap index stocks than on small caps with lower passive ownership.

The result: demographic outflows do not hit "the market" uniformly. They concentrate in index-heavy large caps, potentially creating: - Disproportionate price declines in

large caps - Relative outperformance of small caps and non-index stocks - Opportunities for investors who tilt away from market-cap-weighted indices

This is a natural discussion point for the capstone and connects to Haddad et al.'s finding that pass-through varies across stocks.

## The Literature Gap Our Capstone Fills

Now we can precisely state the capstone's novel contribution. Three literatures exist independently:

**Literature 1: Demographic asset meltdown debate** (Abel, Poterba, Geanakoplos) — asks whether aging populations depress asset prices. Finds mixed results using aggregate approaches in the pre-passive era.

**Literature 2: Passive investing and market structure** (Gabaix-Koijen, Haddad et al., Chinco-Sammon, Ben-David et al.) — shows that passive investing makes markets inelastic and that flows have amplified price effects. Does not apply these findings to demographic flows.

**Literature 3: Retirement savings behavior** (Madrian-Shea, Parker et al., Poterba-Venti-Wise, Mortenson et al.) — documents how people actually save and withdraw in retirement accounts. Does not connect to market-level price impacts.

**The gap:** No existing paper takes concrete, calibrated parameters from Literature 3, feeds them through the institutional channels documented in Literature 3, and applies the price-impact multipliers from Literature 2 to produce a quantitative projection of market effects from demographic flows.

The capstone does exactly this: an accounting model that combines cohort demographics (Lit 1), institutional flow mechanics (Lit 3), and modern price-impact estimates (Lit 2) into a single integrated projection.

## Assumptions and Limitations

An honest treatment of limitations is essential for credibility with examiners. Here are the main ones:

**1. Multiplier symmetry.** Gabaix-Koijen's multiplier was estimated during a period of predominantly *inflows*. Whether it applies symmetrically to sustained *outflows* is unknown. If there is a nonlinearity — if outflows trigger panic selling that magnifies

the effect, or if corporate/government responses dampen it — the multiplier could be different in the outflow regime. We bracket this uncertainty with scenarios.

**2. Partial equilibrium.** Our model is an accounting model, not a general equilibrium model. It does not endogenize how prices, corporate behavior, government policy, or saving rates respond to the projected outflows. Abel (2001) reminds us that general equilibrium effects can partially offset partial equilibrium predictions. Our scenarios attempt to bracket this.

**3. Behavioral stability.** We assume participation rates, contribution rates, and withdrawal rates remain roughly at current levels. Major legislative changes, economic crises, or behavioral shifts could alter these parameters.

**4. Cross-sectional vs. aggregate elasticity.** Haddad et al.'s estimates are for individual stock elasticities; Gabaix-Koijen's are for the aggregate market. We use them together with appropriate caveats, but they are not directly combinable in magnitudes.

**5. U.S. domestic flows only.** The model covers U.S. domestic DC plan flows. It ignores IRA distributions, DB plan runoffs, taxable account liquidations, and international capital flows. This means our projection is a lower bound (addressing only the institutional channel), not a comprehensive forecast.

**6. Buyback uncertainty.** Corporate buybacks are the largest counterforce and are fundamentally unpredictable over a 20-year horizon.

> ***Key Takeaway:*** *These limitations are honest, clearly stated, and addressable through scenario analysis. The capstone does not claim to predict the future — it claims to provide a rigorously grounded projection of one specific, important channel (passive DC plan flows) using the best available academic estimates.*

## The Framing Sentence

Here is the intellectual core of the capstone, distilled into one sentence:

**"The same mechanism that boosted passive equity returns for two decades — steady inflows amplified by an approximately 5× multiplier in an inelastic market — will drag on them for the next two decades as demographic flows reverse."**

This sentence connects the past (inflows → price increases during boomer saving phase), the mechanism (the multiplier), the present (flow reversal as boomers retire), and the future (projected price-level effects). It is the sentence the entire capstone is built to justify and quantify.

## Chapter 7 Summary

**Core concepts:** The capstone fills a specific gap at the intersection of three literatures. The twelve papers provide the theoretical foundation (demographics move prices via life-cycle saving), the institutional parameters (auto-enrollment, TDFs, RMDs, withdrawal behavior), and the price-impact methodology (the 5× multiplier, time-varying with passive growth). Cascading problems of inelasticity extend the implications beyond portfolio values. Buybacks are the key counterforce but are unreliable. International flows reinforce rather than offset domestic pressure.

**The five modules mapped to the literature:** - Module 1 (Demographics): Census data + Geanakoplos MY ratio logic - Module 2 (Inflows): Madrian & Shea (participation) + Parker et al. (TDF glide paths) - Module 3 (Outflows): RMD tables + SECURE 2.0 + Poterba/Venti/Wise (withdrawal data) - Module 4 (Net Flow): Accounting aggregation with Monte Carlo uncertainty - Module 5 (Price Impact): Gabaix-Koijen (multiplier level) + Haddad et al. (multiplier trend)

**Examiner defense:** The capstone's strength is that it does not depend on any single paper or methodology. It builds on converging evidence from independent research teams. The limitations are clearly stated and addressed through scenario analysis.

---

*End of teaching document.*

*Written for Antonio — IE University, BDBA Capstone, February 2026*

*This document synthesizes twelve academic papers across three literatures to build the theoretical foundation for the capstone model. It is a personal study resource designed for comprehensive understanding, not a formal submission.*

---